

The Art of Possible



CAKE SOLUTIONS
ENTERPRISE SOFTWARE SERVICES

MANCHESTER

LONDON

NEW YORK

Chaos Engineering and Runtime Monitoring of Distributed Reactive Systems

Dr. Carl Pulley
Cake Solutions @cakesolutions



The Goal



CAKE SOLUTIONS

The Goal

- Verify the correctness of a distributed system

The Goal

- Verify the correctness of a distributed system
- Typical verification approaches
 - Testing
 - Static analysis
 - Formal/model based verification

What Happens Next?



CAKE SOLUTIONS

What Happens Next?

- Excellent, we now feel confident that the application is correct

What Happens Next?

- Excellent, we now feel confident that the application is correct
- But what happens next?

What About Correctness?



CAKE SOLUTIONS

What About Correctness?

- You Docker containerise your code



What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code
- Human intervention may occur during crisis scenarios

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code
- Human intervention may occur during crisis scenarios
- Do you still feel confident that the application is correct?

The Problem



CAKE SOLUTIONS

The Problem

- Deploy our proven correct application code into dynamic environments



The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes
- Potential for unanticipated failure is now highly likely

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes
- Potential for unanticipated failure is now highly likely
 - embrace and test for failure!

docker-compose-testkit

docker-compose-testkit

- Open-source Scala library (under development)
 - (optionally) deploys, orchestrates and instruments Docker container code
 - agnostic of deployment environment
 - composable behaviour properties
 - reusable Chaos experiments *easily* defined

Demo Recipe



CAKE SOLUTIONS

Demo Recipe

- (optional) describe application containers and networking



Demo Recipe

- (optional) describe application containers and networking
- (partially optional) define system instrumentation
 - view system as a black box
 - at the very least, need to define a layer of interaction code here!



Demo Recipe

- (optional) describe application containers and networking
- (partially optional) define system instrumentation
 - view system as a black box
 - at the very least, need to define a layer of interaction code here!
- define the Chaos experiment
 - relate observed behaviours to data injection

Demo Recipe

- (optional) describe application containers and networking
- (partially optional) define system instrumentation
 - view system as a black box
 - at the very least, need to define a layer of interaction code here!
- define the Chaos experiment
 - relate observed behaviours to data injection
- define runtime monitors (c.f. behaviours)
 - bind monitors to instrumented interface

Demo Overview



CAKE SOLUTIONS

Demo Overview

Networks

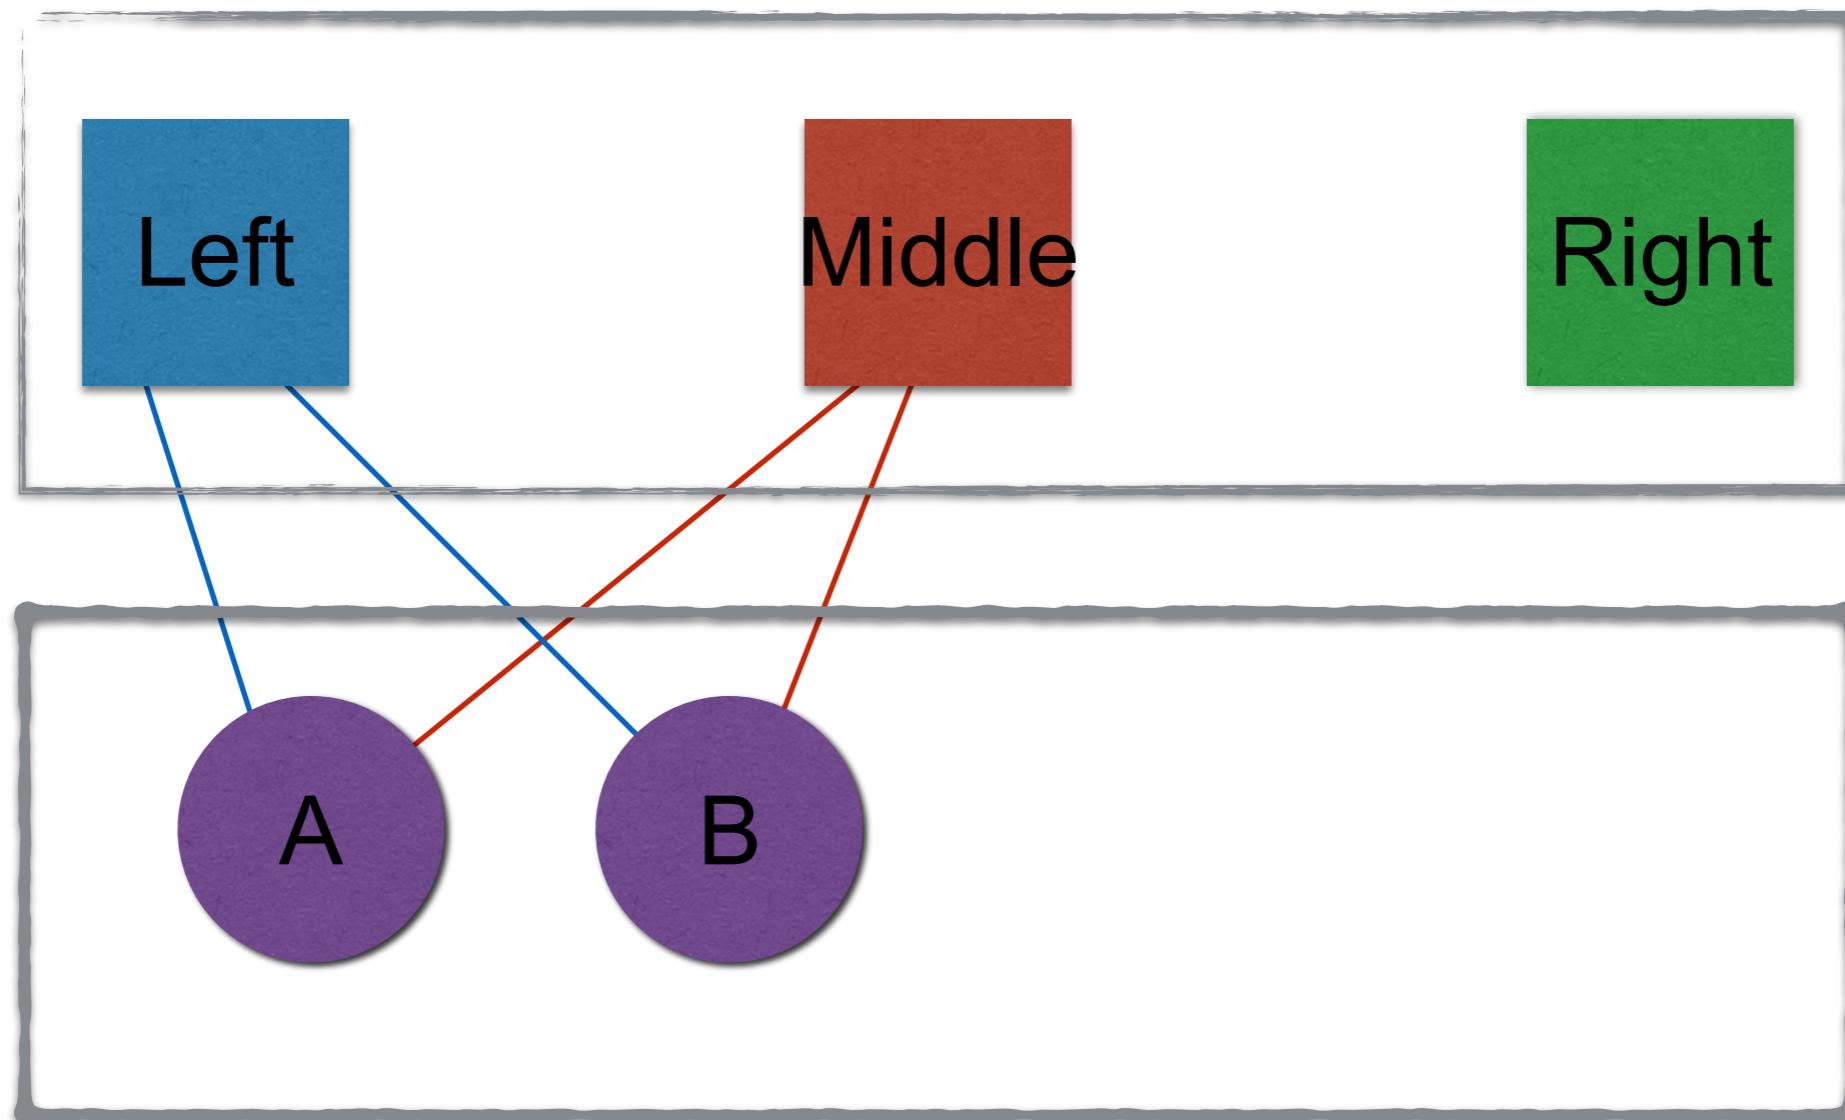
Left

Middle

Right

Demo Overview

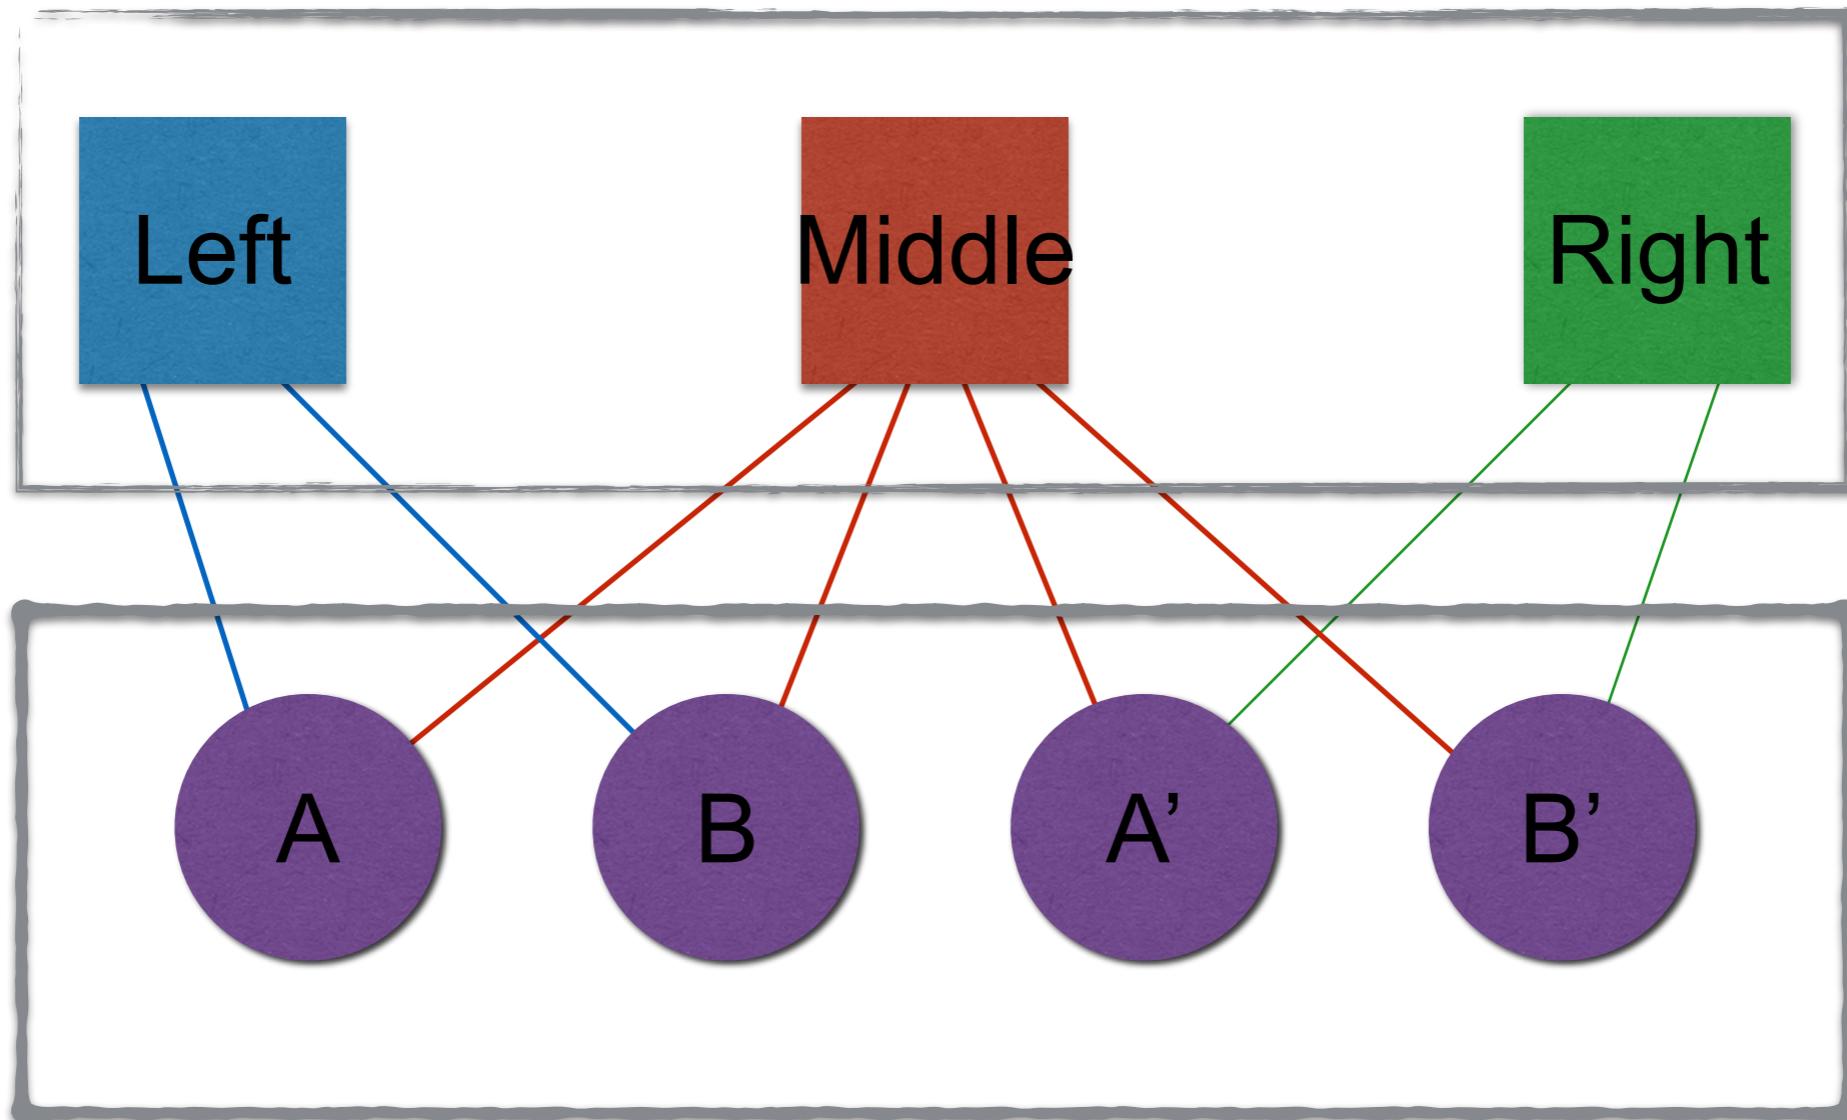
Networks



Akka cluster

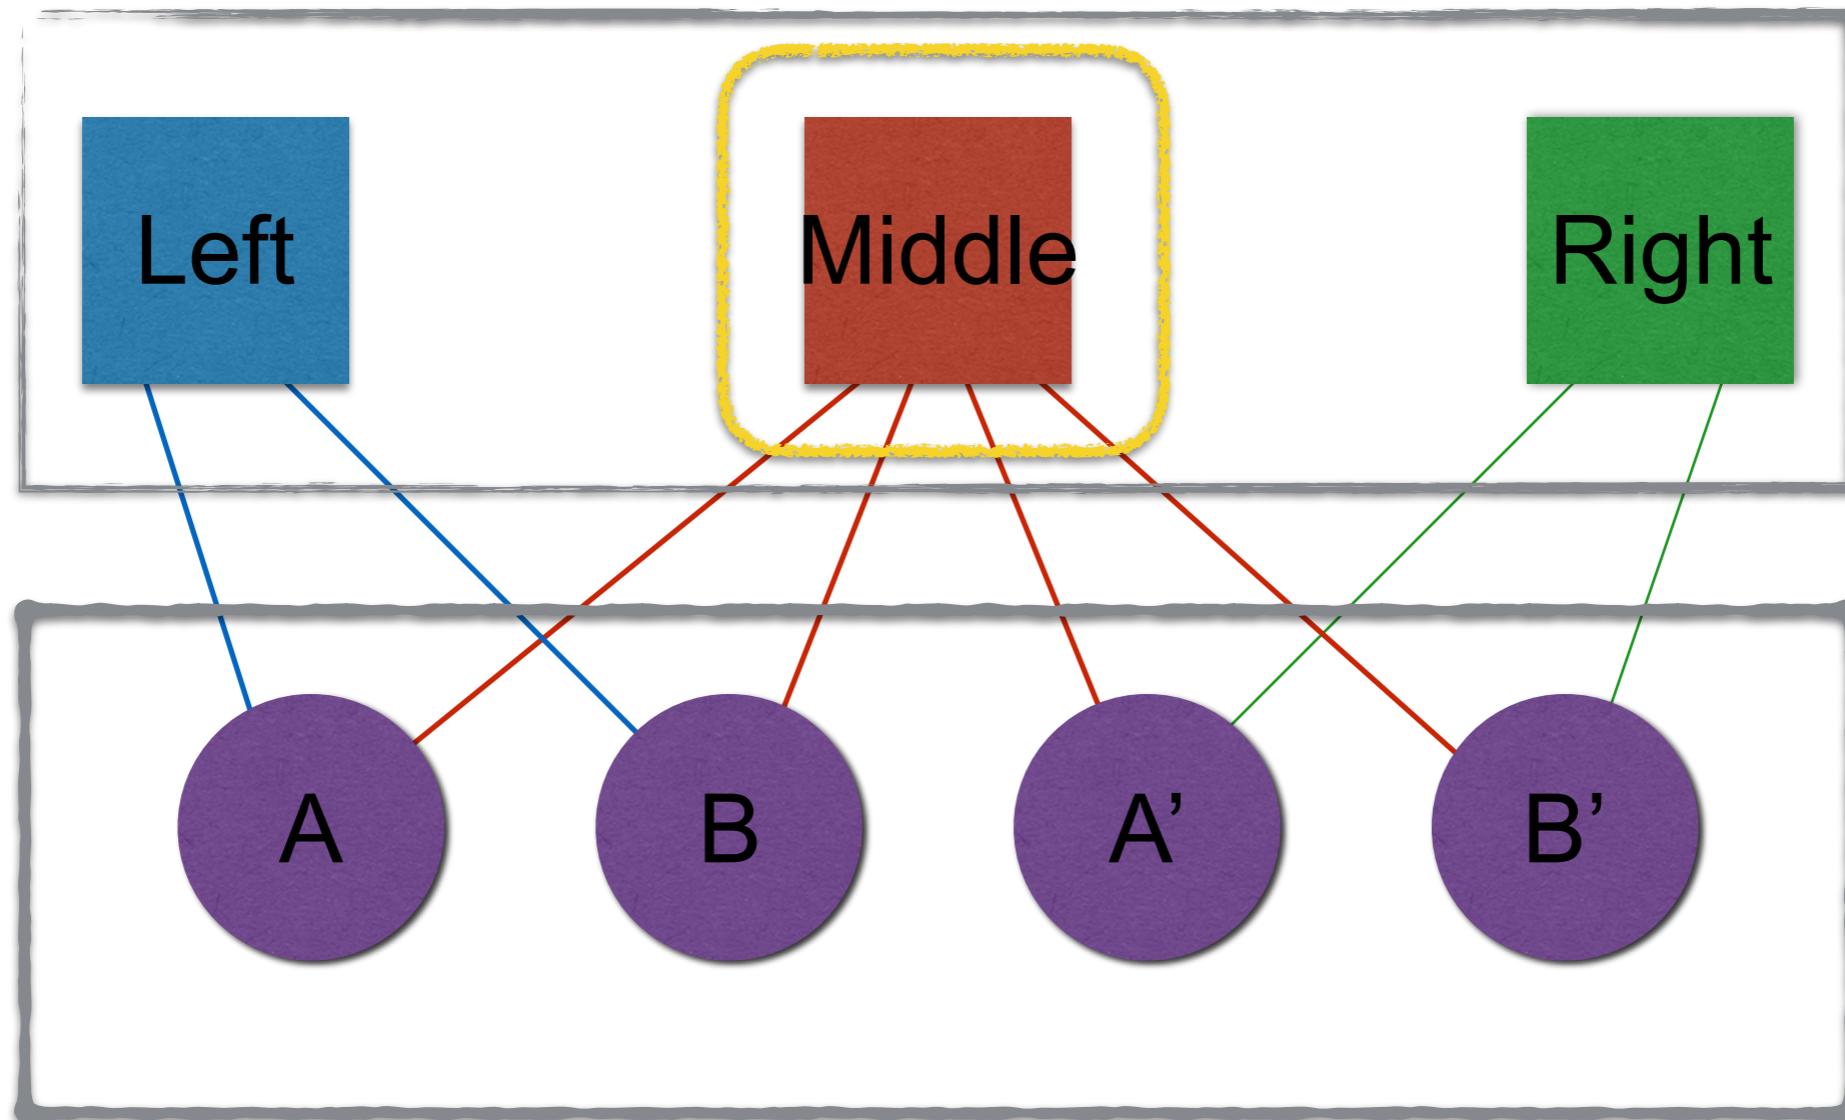
Demo Overview

Networks



Demo Overview

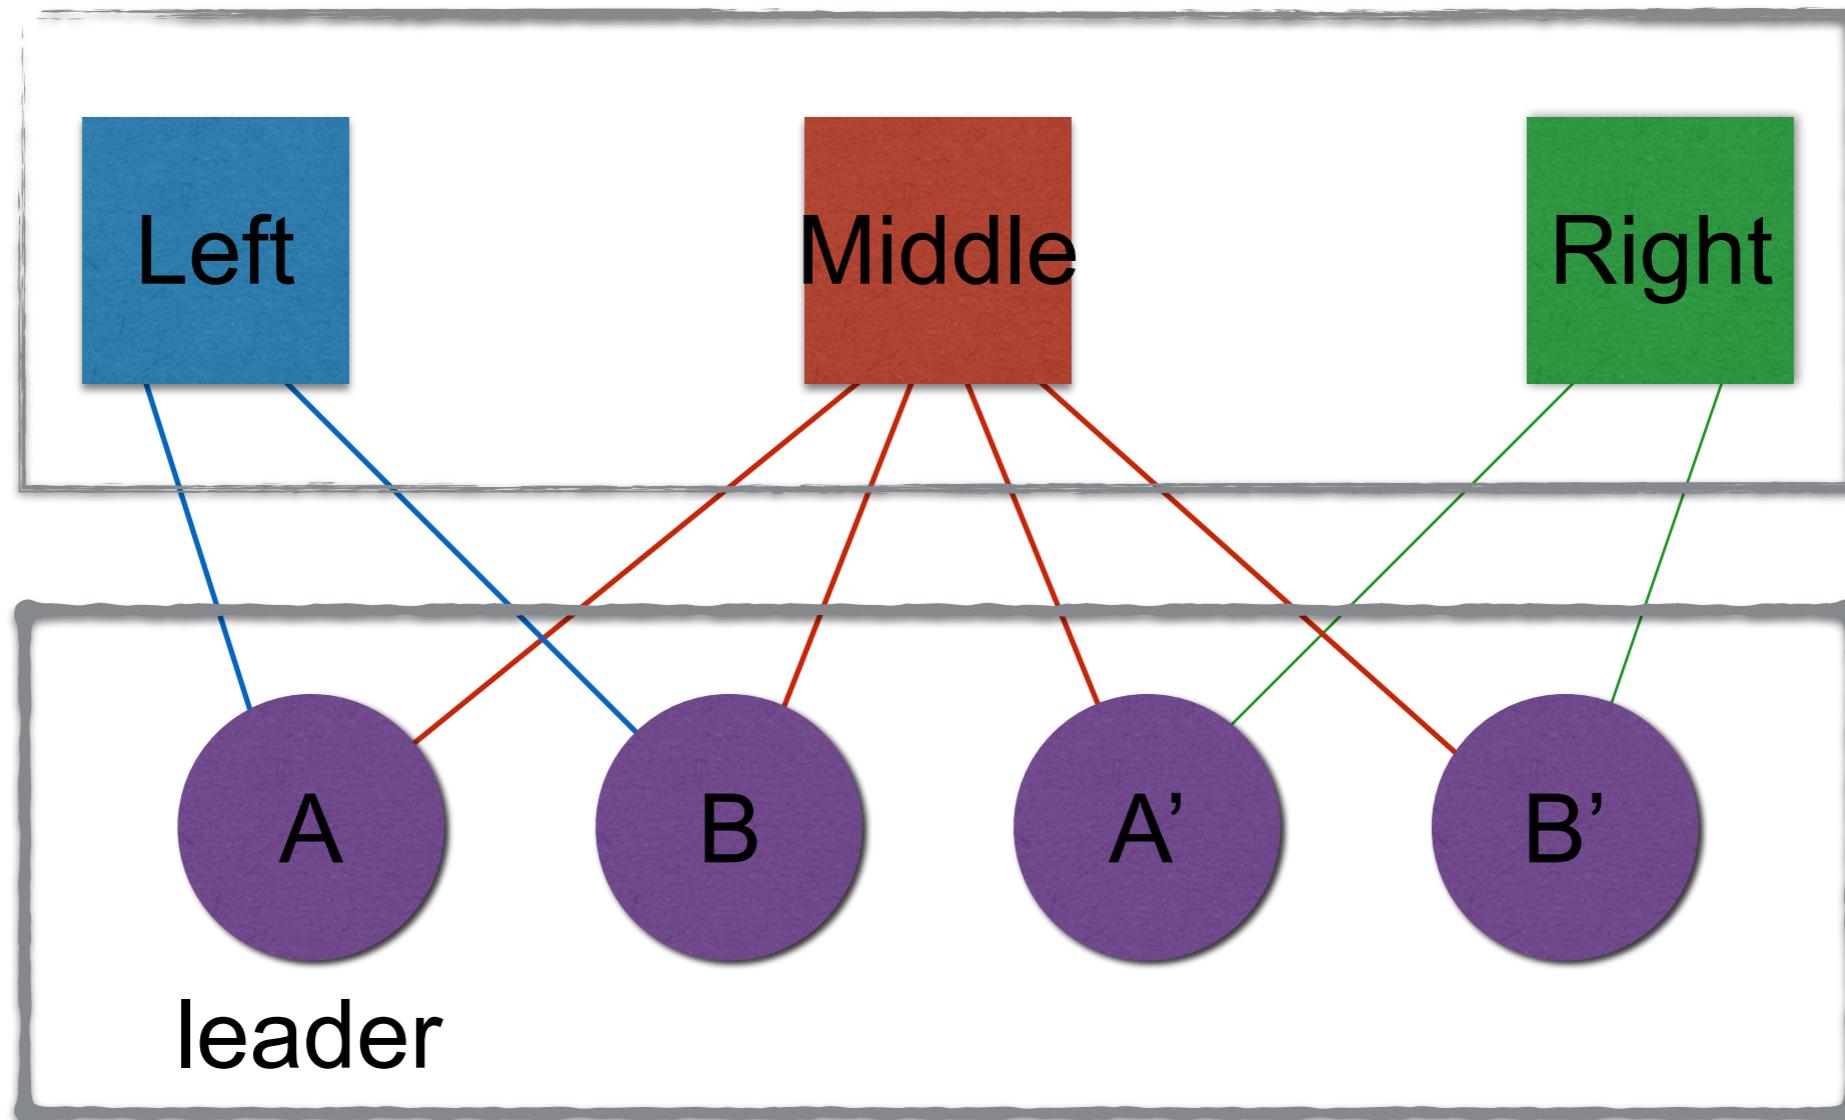
Networks



Akka cluster

Demo Overview

Networks



Auto-downing Enabled!!

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    | template:  
    |   resources:  
    |     - cakesolutions.docker.jmx.akka  
    |     - cakesolutions.docker.network.default.linux  
    |   image: docker-compose-testkit-tests:$version  
    | environment:  
    |   AKKA_HOST: $name  
    |   AKKA_PORT: 2552  
    |   CLUSTER_SEED_NODE: "akka.tcp://TestCluster@$A:2552"  
    | expose:  
    |   - 2552  
    | networks:  
    |   - $network1  
    |   - $network2  
  """.stripMargin  
  
val yaml = DockerComposeString(  
  s"""version: '2'  
    | services:  
    |   ${clusterNode("A", "left", "middle")}  
    |   ${clusterNode("B", "left", "middle")}  
    |   ${clusterNode("A'", "right", "middle")}  
    |   ${clusterNode("B'", "right", "middle")}  
    |  
    | networks:  
    |   left:  
    |   middle:  
    |   right:  
  """.stripMargin  
)
```

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    | template:  
    |   resources:  
    |     - cakesolutions.docker.jmx.akka  
    |     - cakesolutions.docker.network.default.linux  
    |   image: docker-compose-testkit-tests:$version  
    | environment:  
    |   AKKA_HOST: $name  
    |   AKKA_PORT: 2552  
    |   CLUSTER_SEED_NODE: "akka.tcp://TestCluster@$A:2552"  
    | expose:  
    |   - 2552  
    | networks:  
    |   - $network1  
    |   - $network2  
  """.stripMargin
```

```
val yaml = DockerComposeString(  
  s"""version: '2'  
    | services:  
    |   ${clusterNode("A", "left", "middle")}  
    |   ${clusterNode("B", "left", "middle")}  
    |   ${clusterNode("A'", "right", "middle")}  
    |   ${clusterNode("B'", "right", "middle")}  
    |  
    | networks:  
    |   left:  
    |   middle:  
    |   right:  
  """.stripMargin  
)
```

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    | template:  
    |   resources:  
    |     - cakesolutions.docker.jmx.akka  
    |     - cakesolutions.docker.network.default.linux  
    |   image: docker-compose-testkit-tests:$version  
    | environment:  
    |   AKKA_HOST: $name  
    |   AKKA_PORT: 2552  
    |   CLUSTER_SEED_NODE: "akka.tcp://TestCluster@$A:2552"  
    | expose:  
    |   - 2552  
    | networks:  
    |   - $network1  
    |   - $network2  
  """.stripMargin
```

```
val yaml = DockerComposeString(  
  s"""version: '2'  
    | services:  
    |   ${clusterNode("A", "left", "middle")}  
    |   ${clusterNode("B", "left", "middle")}  
    |   ${clusterNode("A'", "right", "middle")}  
    |   ${clusterNode("B'", "right", "middle")}  
    |  
    | networks:  
    |   left:  
    |   middle:  
    |   right:  
  """.stripMargin  
)
```

Demo Overview

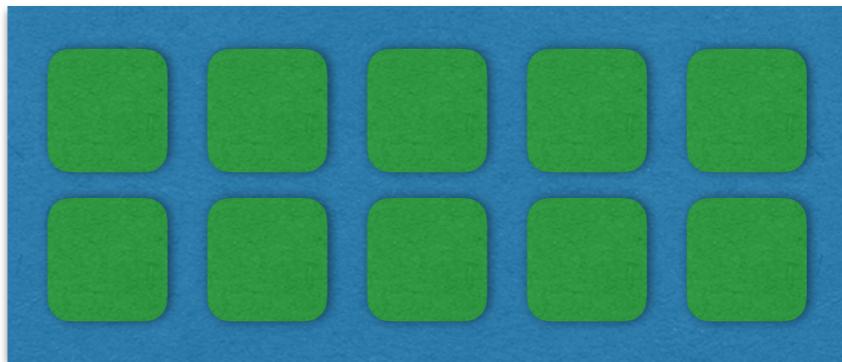
```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    | template:  
    |   resources:  
    |     - cakesolutions.docker.jmx.akka  
    |     - cakesolutions.docker.network.default.linux  
    |   image: docker-compose-testkit-tests:$version  
    |   environment:  
    |     AKKA_HOST: $name  
    |     AKKA_PORT: 2552  
    |     CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"  
    |   expose:  
    |     - 2552  
    |   networks:  
    |     - $network1  
    |     - $network2  
    """.stripMargin  
  
val yaml = DockerComposeString(  
  s"""version: '2'  
    | services:  
    |   ${clusterNode("A", "left", "middle")}  
    |   ${clusterNode("B", "left", "middle")}  
    |   ${clusterNode("A'", "right", "middle")}  
    |   ${clusterNode("B'", "right", "middle")}  
    |  
    | networks:  
    |   left:  
    |   middle:  
    |   right:  
    """.stripMargin  
)
```

Demo Overview

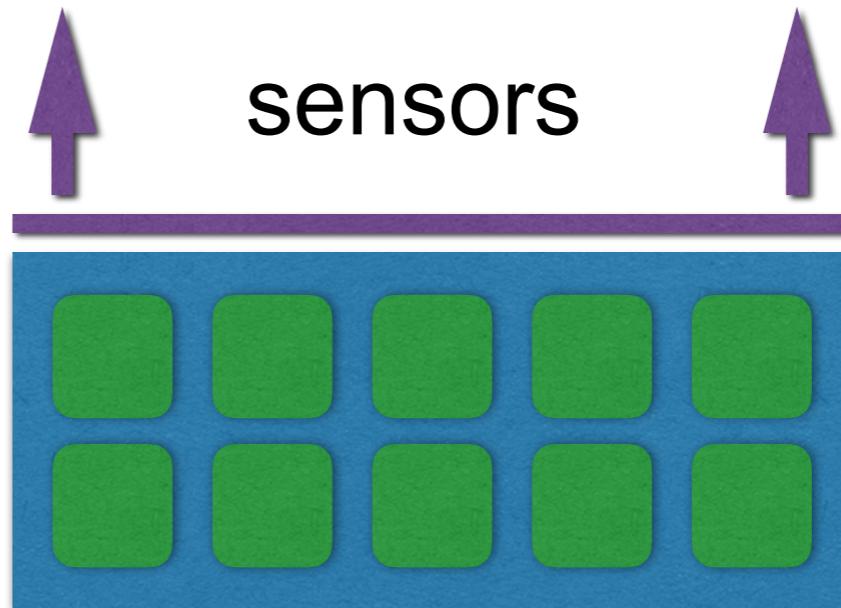
```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    | template:  
    |   resources:  
    |     - cakesolutions.docker.jmx.akka  
    |     - cakesolutions.docker.network.default.linux  
    |   image: docker-compose-testkit-tests:$version  
    |   environment:  
    |     AKKA_HOST: $name  
    |     AKKA_PORT: 2552  
    |     CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"  
    |   expose:  
    |     - 2552  
    |   networks:  
    |     - $network1  
    |     - $network2  
    """.stripMargin  
  
val yaml = DockerComposeString(  
  s"""version: '2'  
    | services:  
    |   ${clusterNode("A", "left", "middle")}  
    |   ${clusterNode("B", "left", "middle")}  
    |   ${clusterNode("A'", "right", "middle")}  
    |   ${clusterNode("B'", "right", "middle")}  
    |  
    | networks:  
    |   left:  
    |   middle:  
    |   right:  
    """.stripMargin  
)
```

immutable container

System Instrumentation

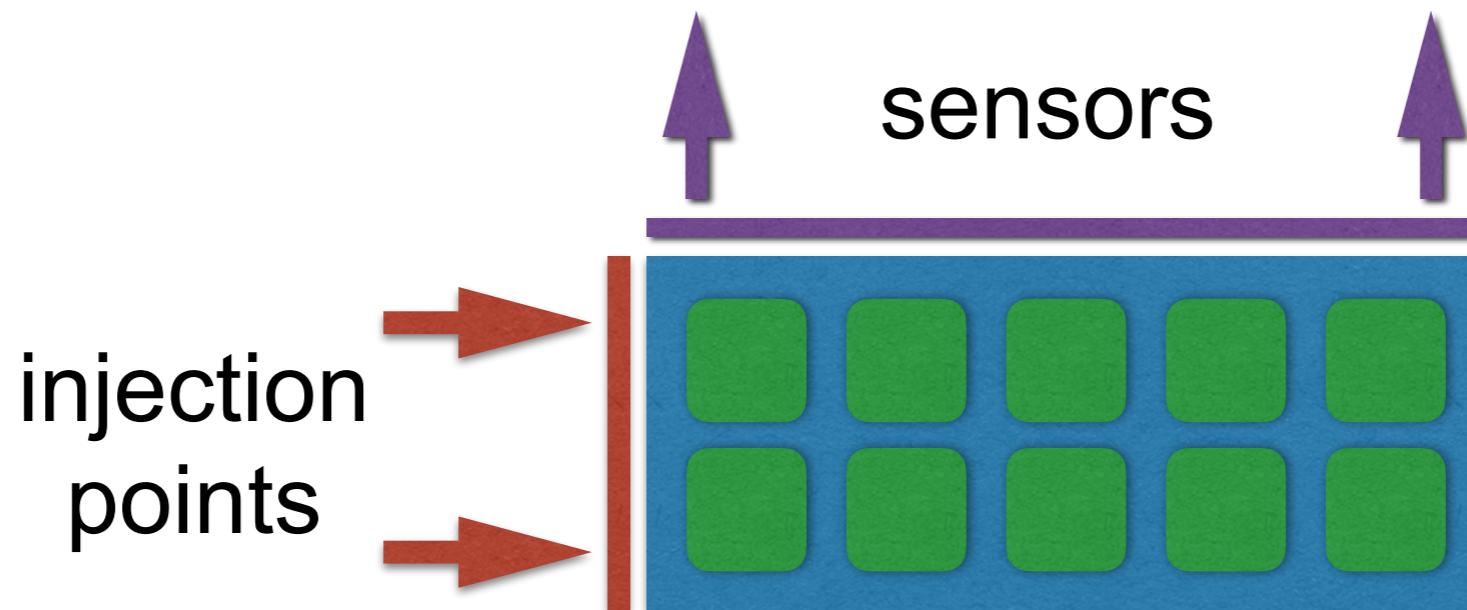


System Instrumentation



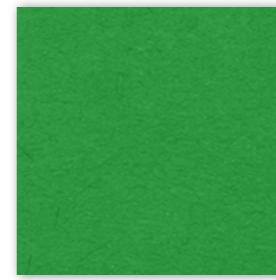
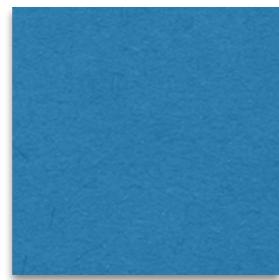
- define sensors
 - emit events of a fixed type
 - implemented using Monix Observables

System Instrumentation



- define sensors
 - emit events of a fixed type
 - implemented using Monix Observables
- define injection points
 - allow typed data to influence the application
 - implemented using Monix Observers

Demo Overview



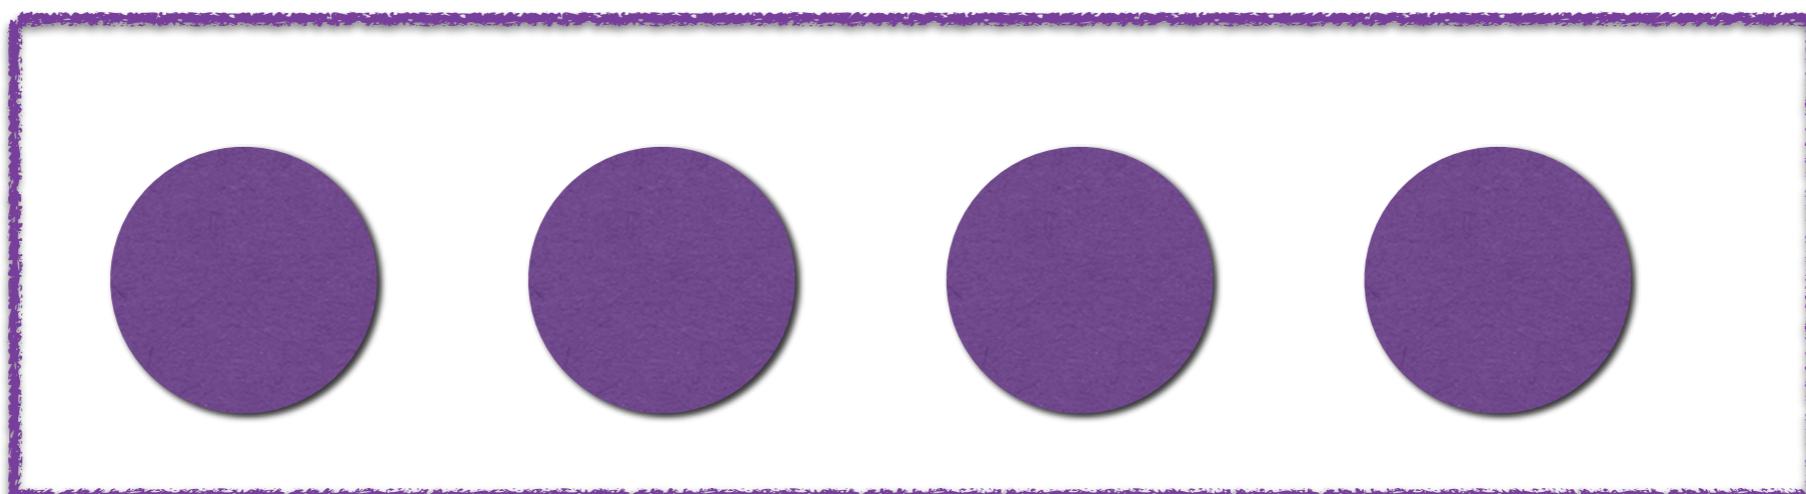
Demo Overview



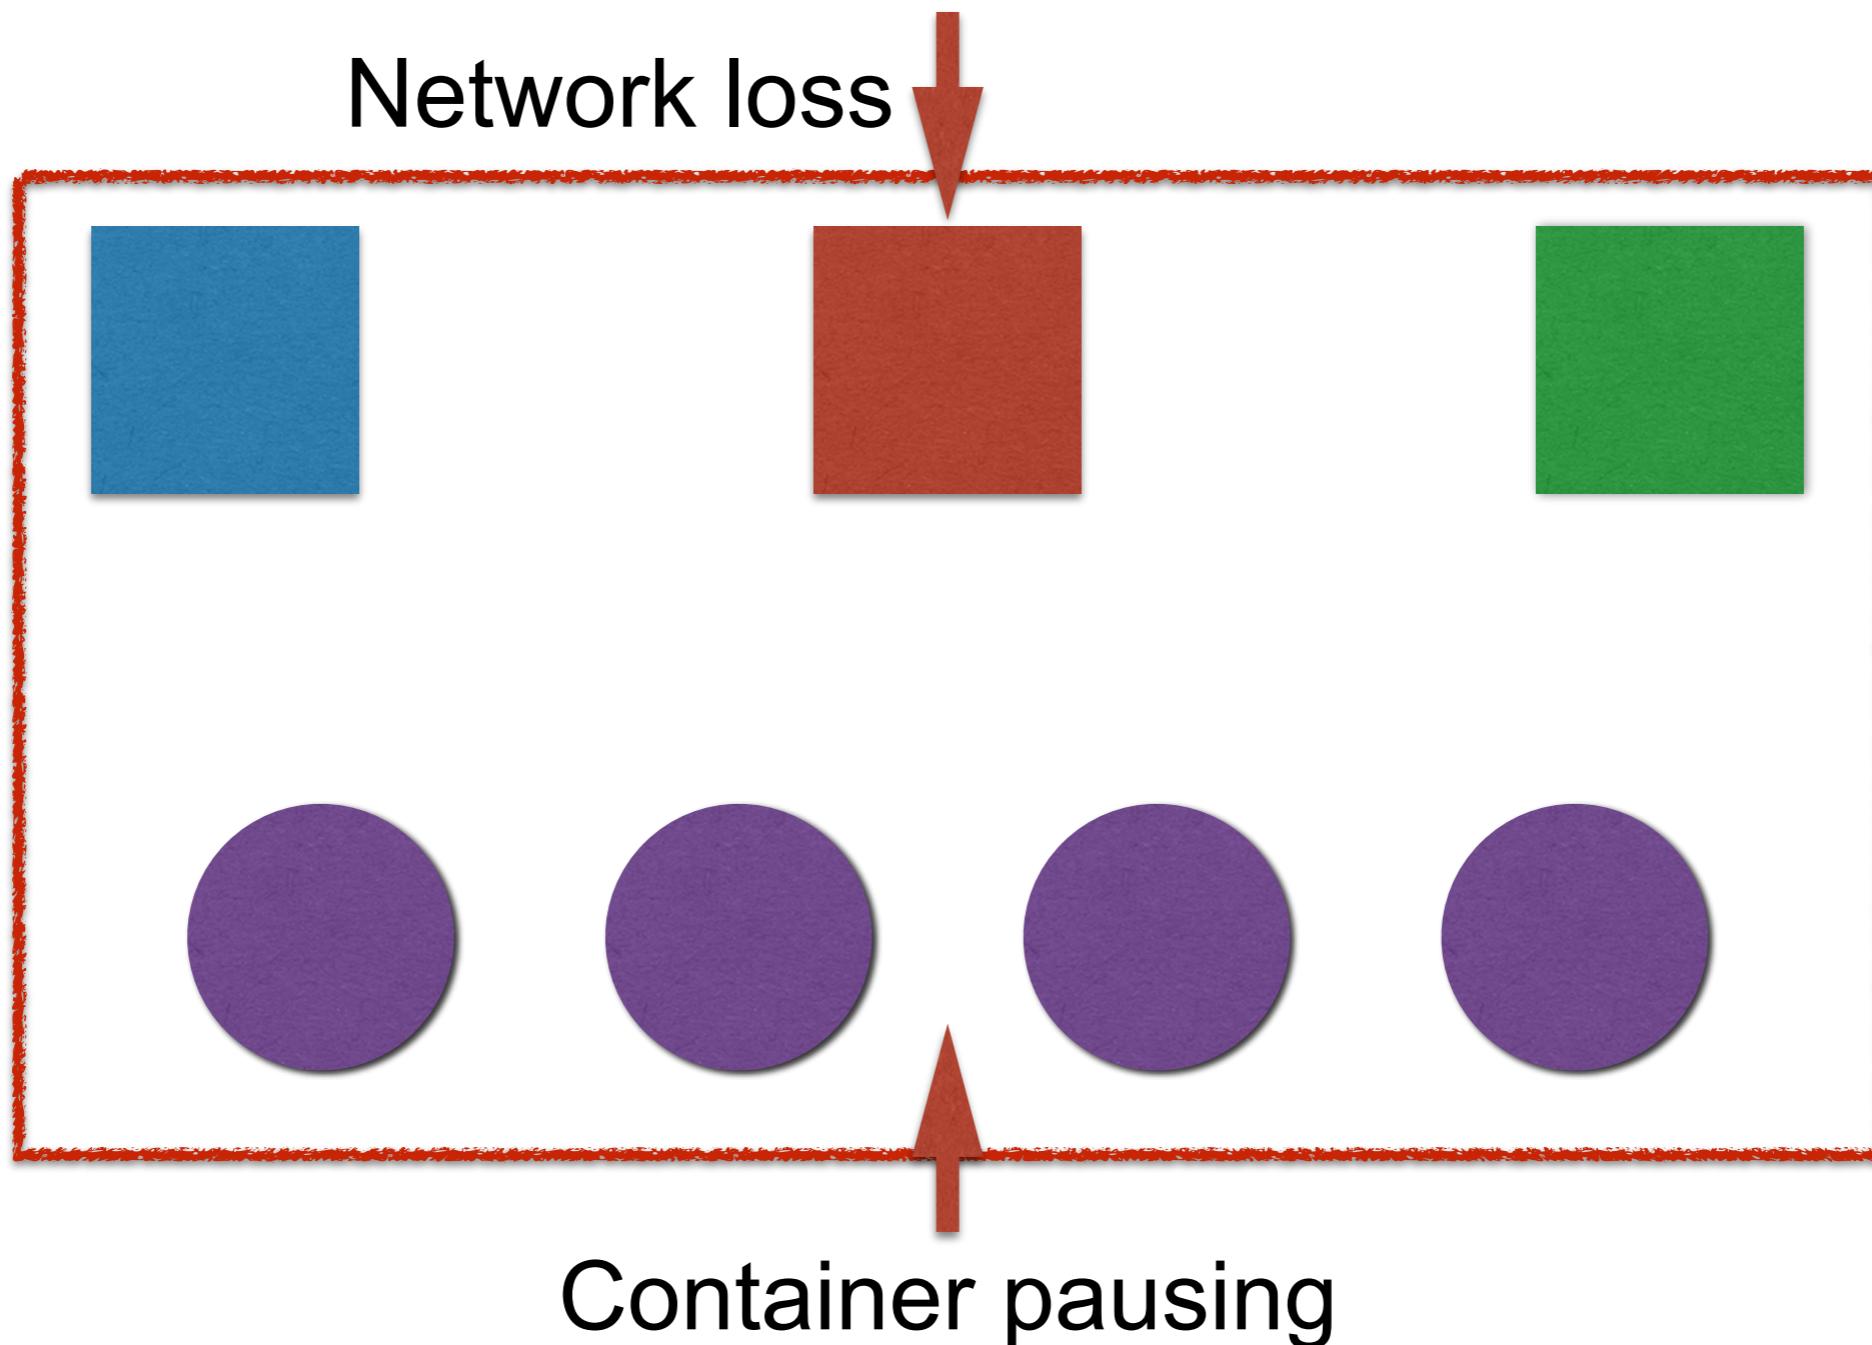
Container
logging



Akka JMX
Console



Demo Overview



Demo Overview

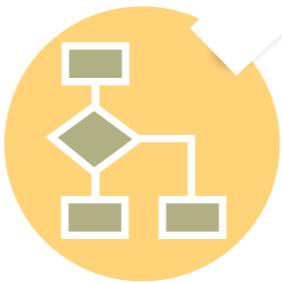
```
final case class AkkaNodeSensors(  
    log: TimedObservable.hot[LogEvent],  
    members: TimedObservable.cold[AkkaClusterState]  
)  
  
object AkkaSensors {  
    def apply(image: DockerImage)(implicit scheduler: Scheduler): AkkaNodeSensors = {  
        AkkaNodeSensors(  
            TimedObservable.hot(image.logging().publish),  
            TimedObservable.cold(image.members()))  
    }  
}  
}
```

Demo Overview

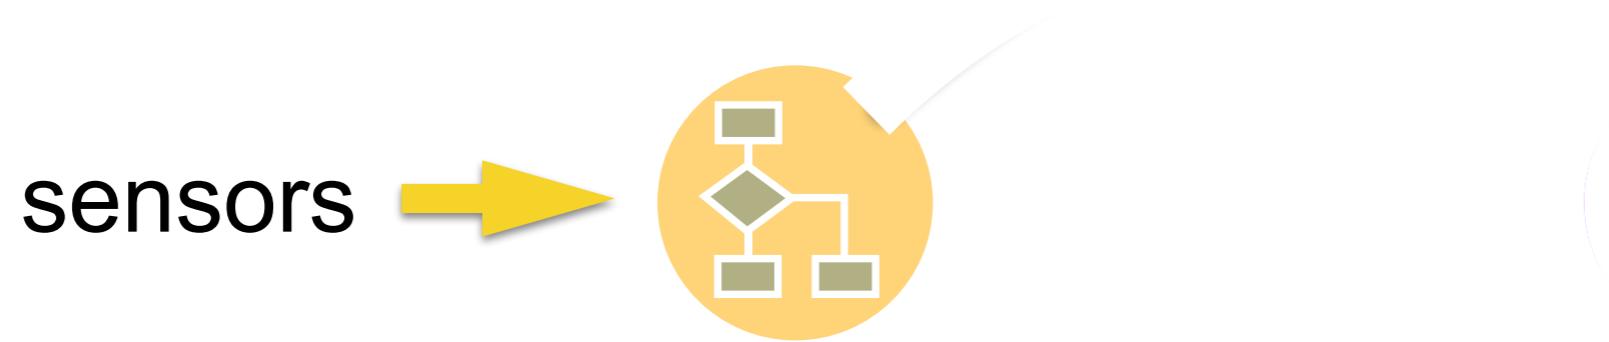
```
final case class AkkaNodeSensors(  
    log: TimedObservable.hot[LogEvent],  
    members: TimedObservable.cold[AkkaClusterState]  
)  
  
object AkkaSensors {  
    def apply(image: DockerImage)(implicit scheduler: Scheduler): AkkaNodeSensors = {  
        AkkaNodeSensors(  
            TimedObservable.hot(image.logging().publish),  
            TimedObservable.cold(image.members()))  
    }  
}
```

For this presentation, injection points will be implemented
using side effecting code

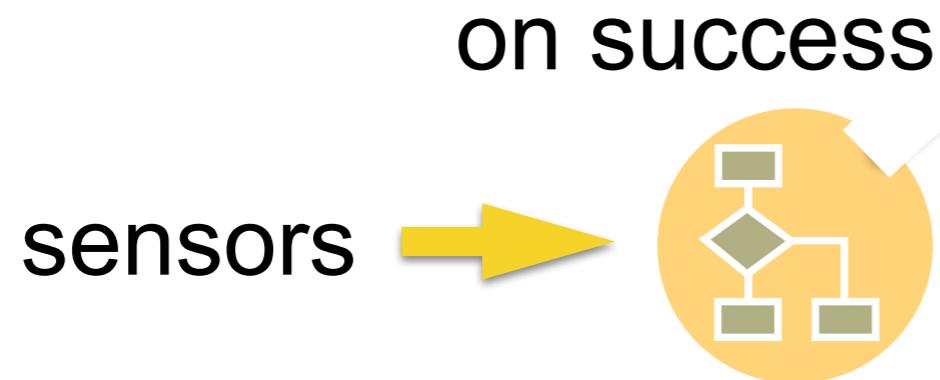
Chaos Experiment



Chaos Experiment

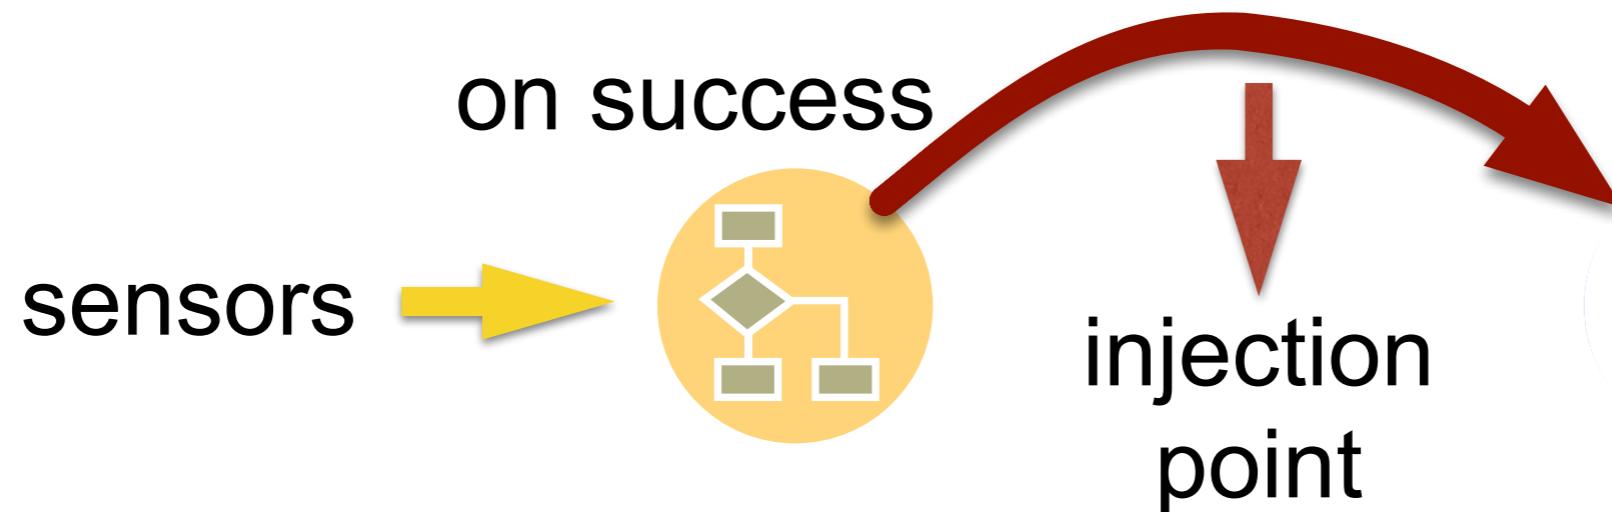


Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour

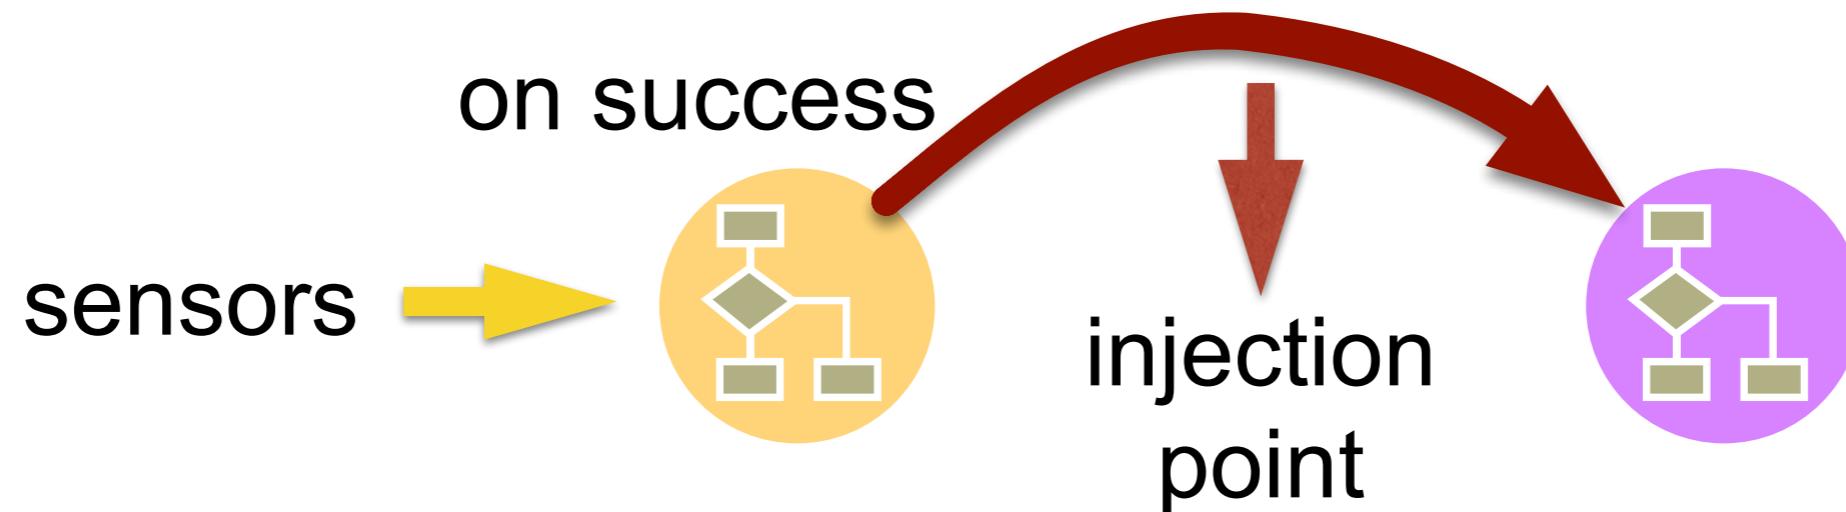
Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour
 - action → happens

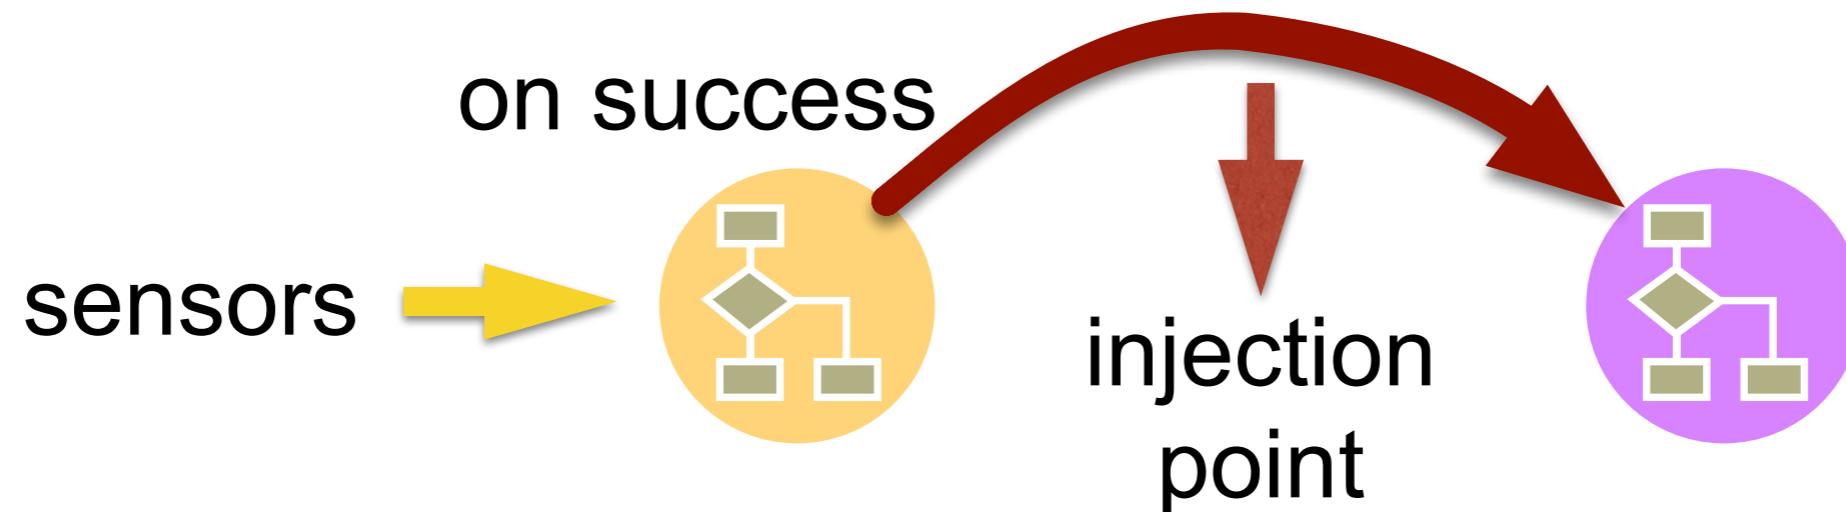


Chaos Experiment



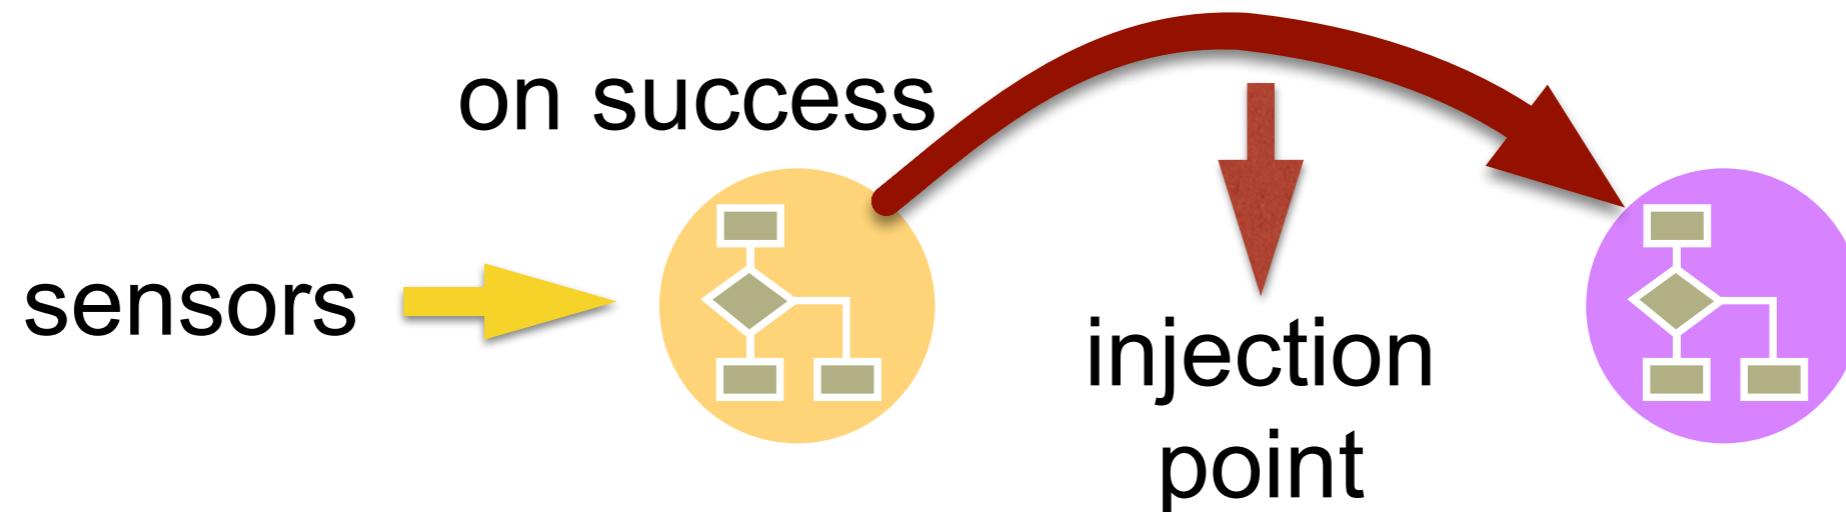
- Interpret this as follows:
 - when we successfully detect behaviour
 - action → happens
 - and we start monitoring for behaviour

Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour
 - action → happens
 - and we start monitoring for behaviour
- Chaining such relations allows the fault injection search space to be described

Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour
 - action → happens
 - and we start monitoring for behaviour
- Chaining such relations allows the fault injection search space to be described

Demo Overview

```
val testSimulation = for {
    _ <- ???
    _ = note("cluster formed, is stable and A an available leader")
    _ = rightNodeA.pause()
    _ = note("A' JVM GC pause starts")
    _ <- ???
    _ = note("A' is unreachable")
    _ = rightNodeA.unpause()
    _ = note("A' JVM GC pause ends")
    _ <- ???
    _ = note("cluster stabilised with A' as a member")
    _ = compose.network("middle").impair(Loss("100%"))
    _ = note("partition into left and right networks")
    _ <- ???
    _ = note("cluster split brains into 3 clusters: A & B; A'; B'"))
} yield Accept()
```

testSimulation should observe(Accept())

Behavioural Property



CAKE SOLUTIONS

Behavioural Property

- State machine



Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state

Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state
 - (potentially) limit how long we may *linger* in a particular state

Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state
 - (potentially) limit how long we may *linger* in a particular state
 - *StateTimeout* is observed when we linger in a given state for too long

Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state
 - (potentially) limit how long we may *linger* in a particular state
 - *StateTimeout* is observed when we linger in a given state for too long
 - may generate notification events

Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state
 - (potentially) limit how long we may *linger* in a particular state
 - *StateTimeout* is observed when we linger in a given state for too long
 - may generate notification events
- Definition consists in defining a series of nested partial functions

Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```

Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```



Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```

Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```



Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```

Demo Overview

```
val stableCluster =  
  MatchingAutomata[StableCluster.type, List[Address]](StableCluster, 3.seconds) {  
  case _ => {  
    case addrs: List[Address @unchecked] if addrs.nonEmpty =>  
      Stop(Fail(s"Detected $addrs as unreachable"))  
    case StateTimeout =>  
      Stop(Accept())  
  }  
}
```

Demo Overview

```
def clusterMembers(nodes: String*) =  
  MatchingAutomata[ClusterMemberCheck.type, AkkaClusterState](ClusterMemberCheck) {  
    case _ => {  
      case AkkaClusterState(_, members, unreachable)  
        if unreachable.isEmpty  
          && members.filter(_.status == Up).flatMap(_.address.host) == Set(nodes: _*) =>  
          Stop(Accept())  
    }  
  }
```



Demo Overview

```
def clusterMembers(nodes: String*) =  
  MatchingAutomata[ClusterMemberCheck.type, AkkaClusterState](ClusterMemberCheck) {  
    case _ => {  
      case AkkaClusterState(_, members, unreachable)  
        if unreachable.isEmpty  
          && members.filter(_.status == Up).flatMap(_.address.host) == Set(nodes: _*) =>  
          Stop(Accept())  
    }  
  }  
  
clusterMembers("A", "B", "A'", "B'").run(clusterSensors("A").members)
```

- Behavioural properties are executed when they bind to 1 or more *physical* sensor streams (i.e. Observables) and injection points (i.e. Observers)

Demo Overview

```
def clusterMembers(nodes: String*) =  
  MatchingAutomata[ClusterMemberCheck.type, AkkaClusterState](ClusterMemberCheck) {  
    case _ => {  
      case AkkaClusterState(_, members, unreachable)  
        if unreachable.isEmpty  
          && members.filter(_.status == Up).flatMap(_.address.host) == Set(nodes: _*) =>  
          Stop(Accept())  
    }  
  }  
  
clusterMembers("A", "B", "A'", "B'").run(clusterSensors("A").members)  
  
available.run(clusterSensors("A").log) && leader.run(clusterSensors("A").log)
```

- Behavioural properties are executed when they bind to 1 or more *physical* sensor streams (i.e. Observables) and injection points (i.e. Observers)
- May combine these using propositional connectives





CAKE SOLUTIONS

Demo Time!



MANCHESTER

LONDON

NEW YORK

Conclusion



CAKE SOLUTIONS

Conclusion

- Seen how to model real-world distributed systems extensionally



Conclusion

- Seen how to model real-world distributed systems extensionally
- Taken steady state behavioural specifications and developed a linear dynamic logic for experiments
 - by relating data injection traces to observed system behaviour

Conclusion

- Seen how to model real-world distributed systems extensionally
- Taken steady state behavioural specifications and developed a linear dynamic logic for experiments
 - by relating data injection traces to observed system behaviour
- Applied techniques to a range of fault-injection scenarios

References

- [https://github.com/carlalley/docker-compose-testkit /tree/ scala-exchange-2016](https://github.com/carlalley/docker-compose-testkit/tree/scala-exchange-2016)
- *Chaos Engineering* by Ali Basiri, Niosha Behnam, Ruud de Rooij, Lorin Hochstein, Luke Kosewski, Justin Reynolds, Casey Rosenthal
- <http://principlesofchaos.org/>
- *The Weird Machines in Proof-Carrying Code* by Julien Vanegue
- *The Byzantine Generals Problem* by L.Lamport, R.Shostak and M.Pease
- *Impossibility of Distributed Consensus with One Faulty Process* by M.Fischer, N.Lynch and M.Paterson



@cakesolutions



0845 617 1200



enquiries@cakesolutions.net



MANCHESTER

LONDON

NEW YORK