



MANCHESTER

LONDON

NEW YORK

Chaos Engineering and Runtime Monitoring of Distributed Reactive Systems

Dr. Carl Pulley

Cake Solutions @cakesolutions



The Goal

- Verify the correctness of a distributed system
- Typical verification approaches
 - Testing
 - Static analysis
 - Formal/model based verification

What Happens Next?

- Excellent, we now feel confident that the application is correct
- But what happens next?

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code
- Human intervention may occur during crisis scenarios
- Do you still feel confident that the application is correct?

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes
- Potential for unanticipated failure is now highly likely
 - embrace and test for failure!


docker-compose-testkit

- Open-source Scala library (under development)
 - (optionally) deploys, orchestrates and instruments Docker container code
 - agnostic of deployment environment
 - composable behaviour properties
 - reusable Chaos experiments *easily* defined

Demo Recipe

- (optional) describe application containers and networking
- (partially optional) define system instrumentation
 - view system as a black box
 - at the very least, need to define a layer of interaction code here!
- define the Chaos experiment
 - relate observed behaviours to data injection
- define runtime monitors (c.f. behaviours)
 - bind monitors to instrumented interface

Demo Overview



Auto-downing Enabled!!

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:
```

```
    template:  
      resources:  
        - cakesolutions.docker.jmx.akka  
        - cakesolutions.docker.network.default.linux  
        image: docker-compose-testkit-tests:$version  
        AKKA_HOST: $name  
        AKKA_PORT: 2552  
        CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"  
      expose:  
        - 2552  
      networks:  
        - $network1  
        - $network2  
  """.stripMargin
```


immutable container

```
val yaml = DockerComposeString(  
  s"""version: '2'
```

```
    services:  
      ${clusterNode("A", "left", "middle")}  
      ${clusterNode("B", "left", "middle")}  
      ${clusterNode("A", "right", "middle")}  
      ${clusterNode("B", "right", "middle")}  
    networks:  
      left:  
      middle:  
      right:  
  """.stripMargin  
)
```





System Instrumentation



- define sensors
 - emit events of a fixed type
 - implemented using Monix Observables
- define injection points
 - allow typed data to influence the application
 - implemented using Monix Observers

Demo Overview


Network loss



Container
logging



Akka JMX
Console



Container pausing

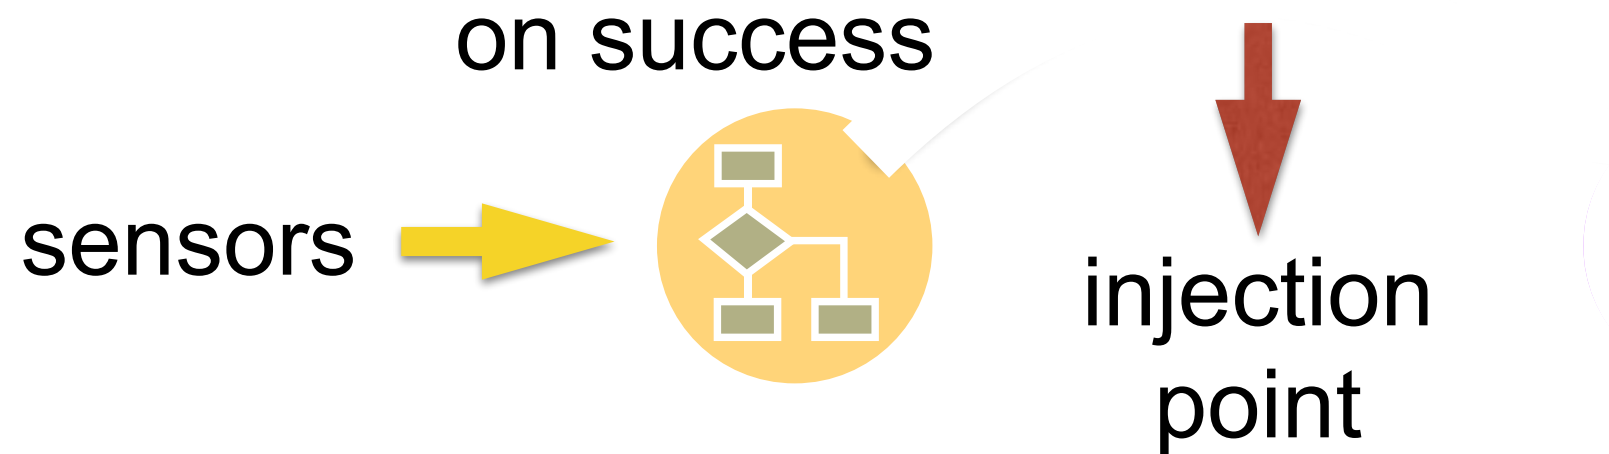


Demo Overview

```
final case class AkkaNodeSensors(  
  log: TimedObservable.hot[LogEvent],  
  members: TimedObservable.cold[AkkaClusterState]  
)  
  
object AkkaSensors {  
  def apply(image: DockerImage)(implicit scheduler: Scheduler): AkkaNodeSensors = {  
    AkkaNodeSensors(  
      TimedObservable.hot(image.logging().publish),  
      TimedObservable.cold(image.members())  
    )  
  }  
}
```

For this presentation, injection points will be implemented using side effecting code

Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour ●
 - action ➡ happens
 - and we start monitoring for behaviour ●
- Chaining such relations allows the fault injection search space to be described

Demo Overview

```
val testSimulation = for {  
  _ <- ???  
  _ = note("cluster formed, is stable and A an available leader")  
  _ = rightNodeA.pause()  
  _ = note("A' JVM GC pause starts")  
  _ <- ???  
  _ = note("A' is unreachable")  
  _ = rightNodeA.unpause()  
  _ = note("A' JVM GC pause ends")  
  _ <- ???  
  _ = note("cluster stabilised with A' as a member")  
  _ = compose.network("middle").impair(Loss("100%"))  
  _ = note("partition into left and right networks")  
  _ <- ???  
  _ = note("cluster split brains into 3 clusters: A & B; A'; B'")  
} yield Accept()
```

```
testSimulation should observe(Accept())
```

Behavioural Property

- State machine
 - sensor events (i.e. Observable) cause machine to change state
 - (potentially) limit how long we may *linger* in a particular state
 - *StateTimeout* is observed when we linger in a given state for too long
 - may generate notification events
- Definition consists in defining a series of nested partial functions

Demo Overview

```
val available =  
  MatchingAutomata[WaitToBeAvailable.type, Boolean](WaitToBeAvailable) {  
    case _ => {  
      case true =>  
        Stop(Accept())  
    }  
  }  
  
val leader =  
  MatchingAutomata[WaitForLeaderElection.type, Address](WaitForLeaderElection) {  
    case _ => {  
      case addr: Address if addr.host.contains("A") =>  
        Stop(Accept())  
    }  
  }
```

Demo Overview

```
val stableCluster =  
  MatchingAutomata[StableCluster.type, List[Address]](StableCluster, 3.seconds) {  
    case _ => {  
      case addr: List[Address @unchecked] if addr.nonEmpty =>  
        Stop(Fail(s"Detected $addr as unreachable"))  
      case StateTimeout =>  
        Stop(Accept())  
    }  
  }
```

Demo Overview

```
def clusterMembers(nodes: String*) =  
  MatchingAutomata[ClusterMemberCheck.type, AkkaClusterState](ClusterMemberCheck) {  
    case _ => {  
      case AkkaClusterState(_, members, unreachable)  
        if unreachable.isEmpty  
          && members.filter(_.status == Up).flatMap(_.address.host) == Set(nodes: _*) =>  
        Stop(Accept())  
    }  
  }  
  
clusterMembers("A", "B", "A'", "B'").run(clusterSensors("A").members)  
  
available.run(clusterSensors("A").log) && leader.run(clusterSensors("A").log)
```

- Behavioural properties are executed when they bind to 1 or more *physical* sensor streams (i.e. Observables) and injection points (i.e. Observers)
- May combine these using propositional connectives



CAKESOLUTIONS

Demo Time!



MANCHESTER

LONDON

NEW YORK

Conclusion

- Seen how to model real-world distributed systems extensionally
- Taken steady state behavioural specifications and developed a linear dynamic logic for experiments
 - by relating data injection traces to observed system behaviour
- Applied techniques to a range of fault-injection scenarios

References

- <https://github.com/carlpulley/docker-compose-testkit/tree/scala-exchange-2016>
- *Chaos Engineering* by Ali Basiri, Niosha Behnam, Ruud de Rooij, Lorin Hochstein, Luke Kosewski, Justin Reynolds, Casey Rosenthal
- <http://principlesofchaos.org/>
- *The Weird Machines in Proof-Carrying Code* by Julien Vanegue
- *The Byzantine Generals Problem* by L.Lamport, R.Shostak and M.Pease
- *Impossibility of Distributed Consensus with One Faulty Process* by M.Fischer, N.Lynch and M.Paterson



CAKESOLUTIONS



@cakesolutions



0845 617 1200



enquiries@cakesolutions.net



MANCHESTER

LONDON

NEW YORK