



MANCHESTER

LONDON

NEW YORK

Chaos Engineering and Runtime Monitoring of Distributed Reactive Systems

Dr. Carl Pulley

Cake Solutions @cakesolutions



Chaos Engineering

Chaos Engineering

- Emphasises an empirical approach to testing and monitoring distributed systems

Chaos Engineering

- Emphasises an empirical approach to testing and monitoring distributed systems
- Typical Chaos Experiment
 - define views of a system (e.g. by monitoring a set of system metrics)
 - randomly inject faults (e.g. by killing containers or changing networking behaviour)

Chaos Engineering

- Emphasises an empirical approach to testing and monitoring distributed systems
- Typical Chaos Experiment
 - define views of a system (e.g. by monitoring a set of system metrics)
 - randomly inject faults (e.g. by killing containers or changing networking behaviour)
- Infer weaknesses using deviations from expected or steady-state behaviour

Chaos Engineering

- Emphasises an empirical approach to testing and monitoring distributed systems
- Typical Chaos Experiment
 - define views of a system (e.g. by monitoring a set of system metrics)
 - randomly inject faults (e.g. by killing containers or changing networking behaviour)
- Infer weaknesses using deviations from expected or steady-state behaviour
- Potentially apply to production systems!

The Goal

The Goal

- Verify the correctness of a distributed system

The Goal

- Verify the correctness of a distributed system
- Typical verification approaches
 - Testing
 - Static analysis
 - Formal/model based verification

What Happens Next?

What Happens Next?

- Excellent, we now feel confident that the application is correct

What Happens Next?

- Excellent, we now feel confident that the application is correct
- But what happens next?

What About Correctness?

What About Correctness?

- You Docker containerise your code

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code
- Human intervention may occur during crisis scenarios

What About Correctness?

- You Docker containerise your code
- DevOps wrap your code up with provisioning code, inject processes into containers, add monitoring and instrumentation code, etc.
- DevOps then deploy to and orchestrate commodity cloud based hardware
- The data centre dynamically maintains/relocates your code
- Human intervention may occur during crisis scenarios
- Do you still feel confident that the application is correct?

The Problem

The Problem

- Deploy our proven correct application code into dynamic environments

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes
- Potential for unanticipated failure is now highly likely

The Problem

- Deploy our proven correct application code into dynamic environments
 - we have little or no real control over these environments!
 - deployment typically uses unverified processes
- Potential for unanticipated failure is now highly likely
 - embrace and test for failure!

docker-compose-testkit

docker-compose-testkit

- Open-source Scala library (under development)
 - (optionally) deploys, orchestrates and instruments Docker container code
 - agnostic of deployment environment
 - composable behaviour properties
 - reusable Chaos experiments *easily* defined

templates

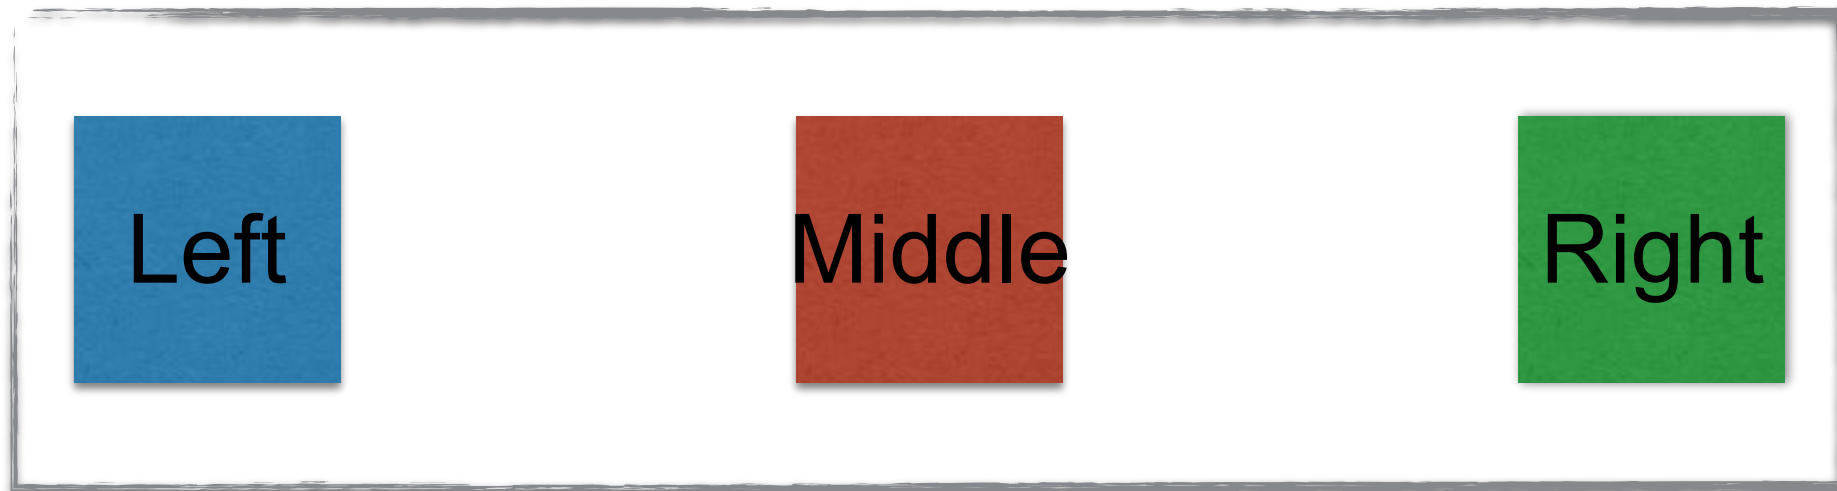
templates

- Specify Docker image instrumentation layers
 - injects code and resources into image
- Specify library code for using instrumentation
 - code needs to use side effects
 - maintain a functional core using extensible effects (i.e. Eff)
 - sensor events represented using Monix Observables

Demo Overview

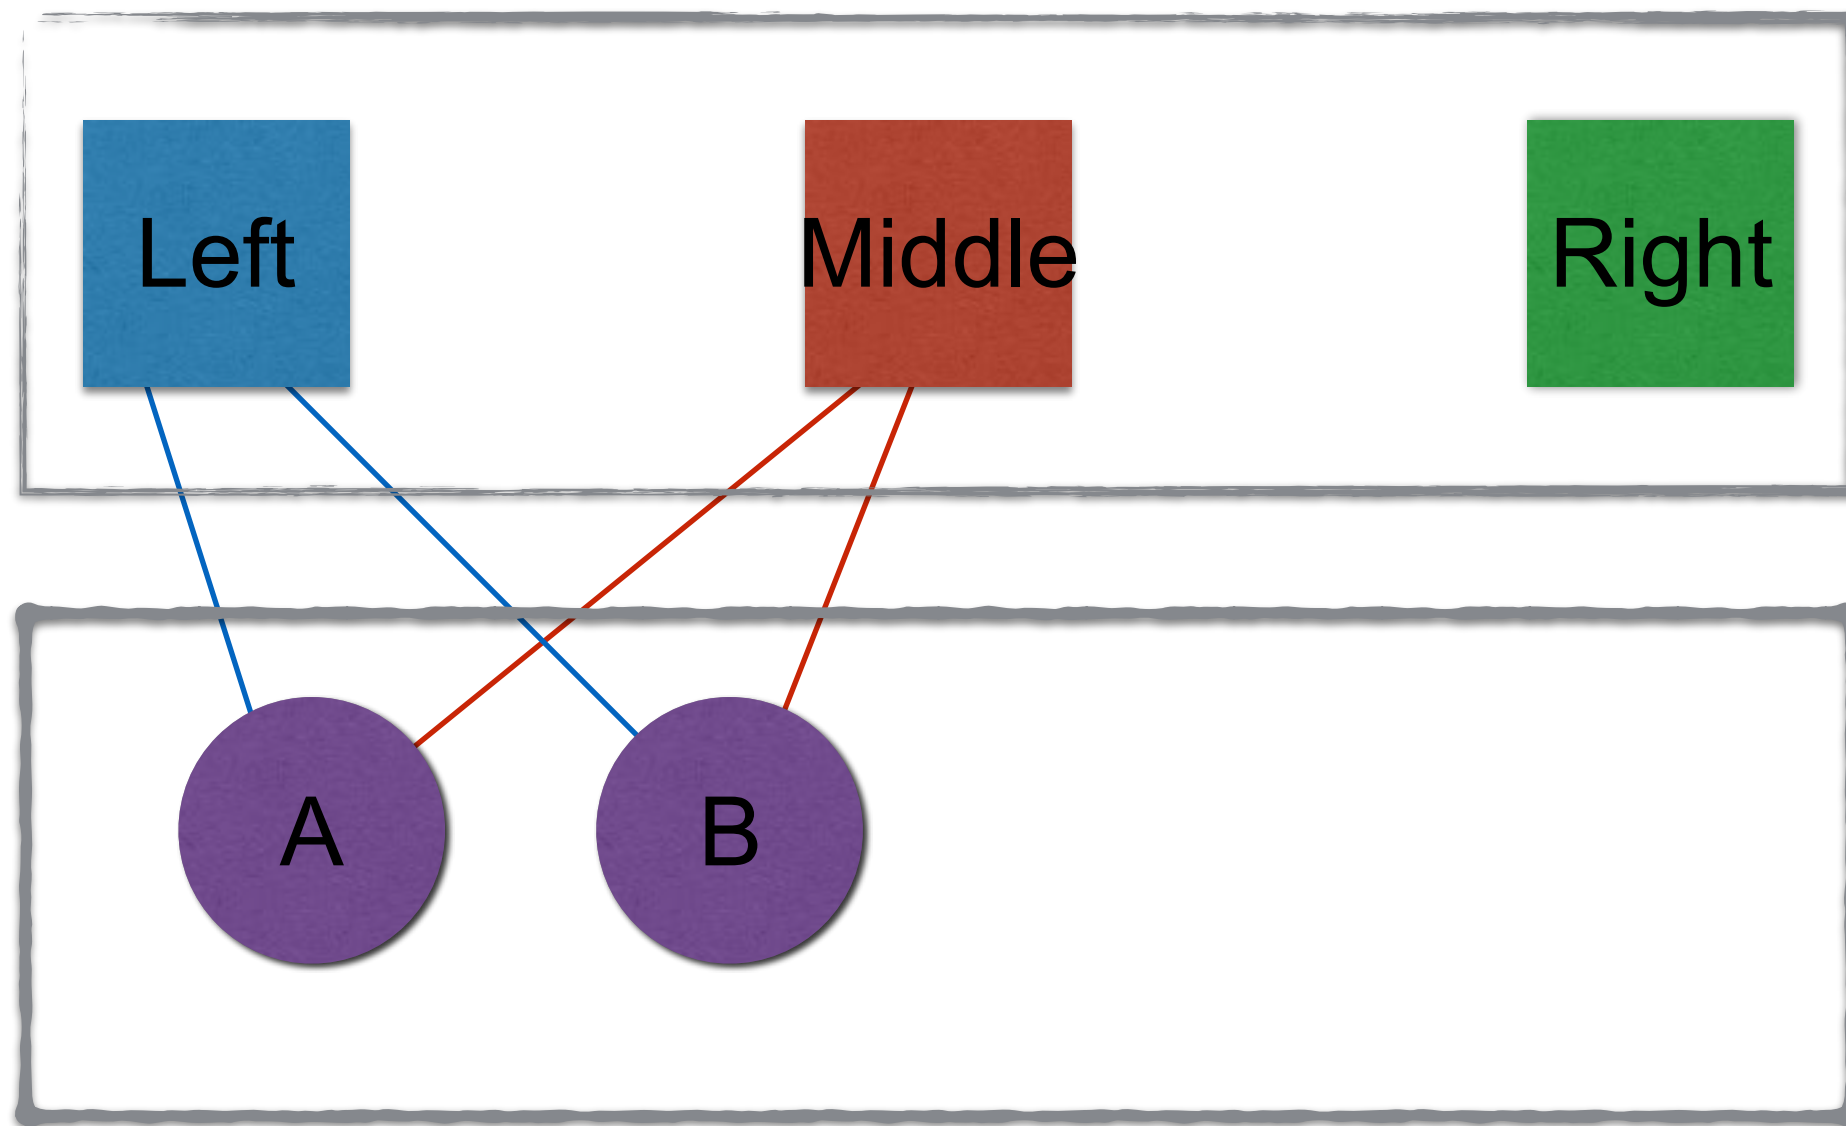
Demo Overview

Networks



Demo Overview

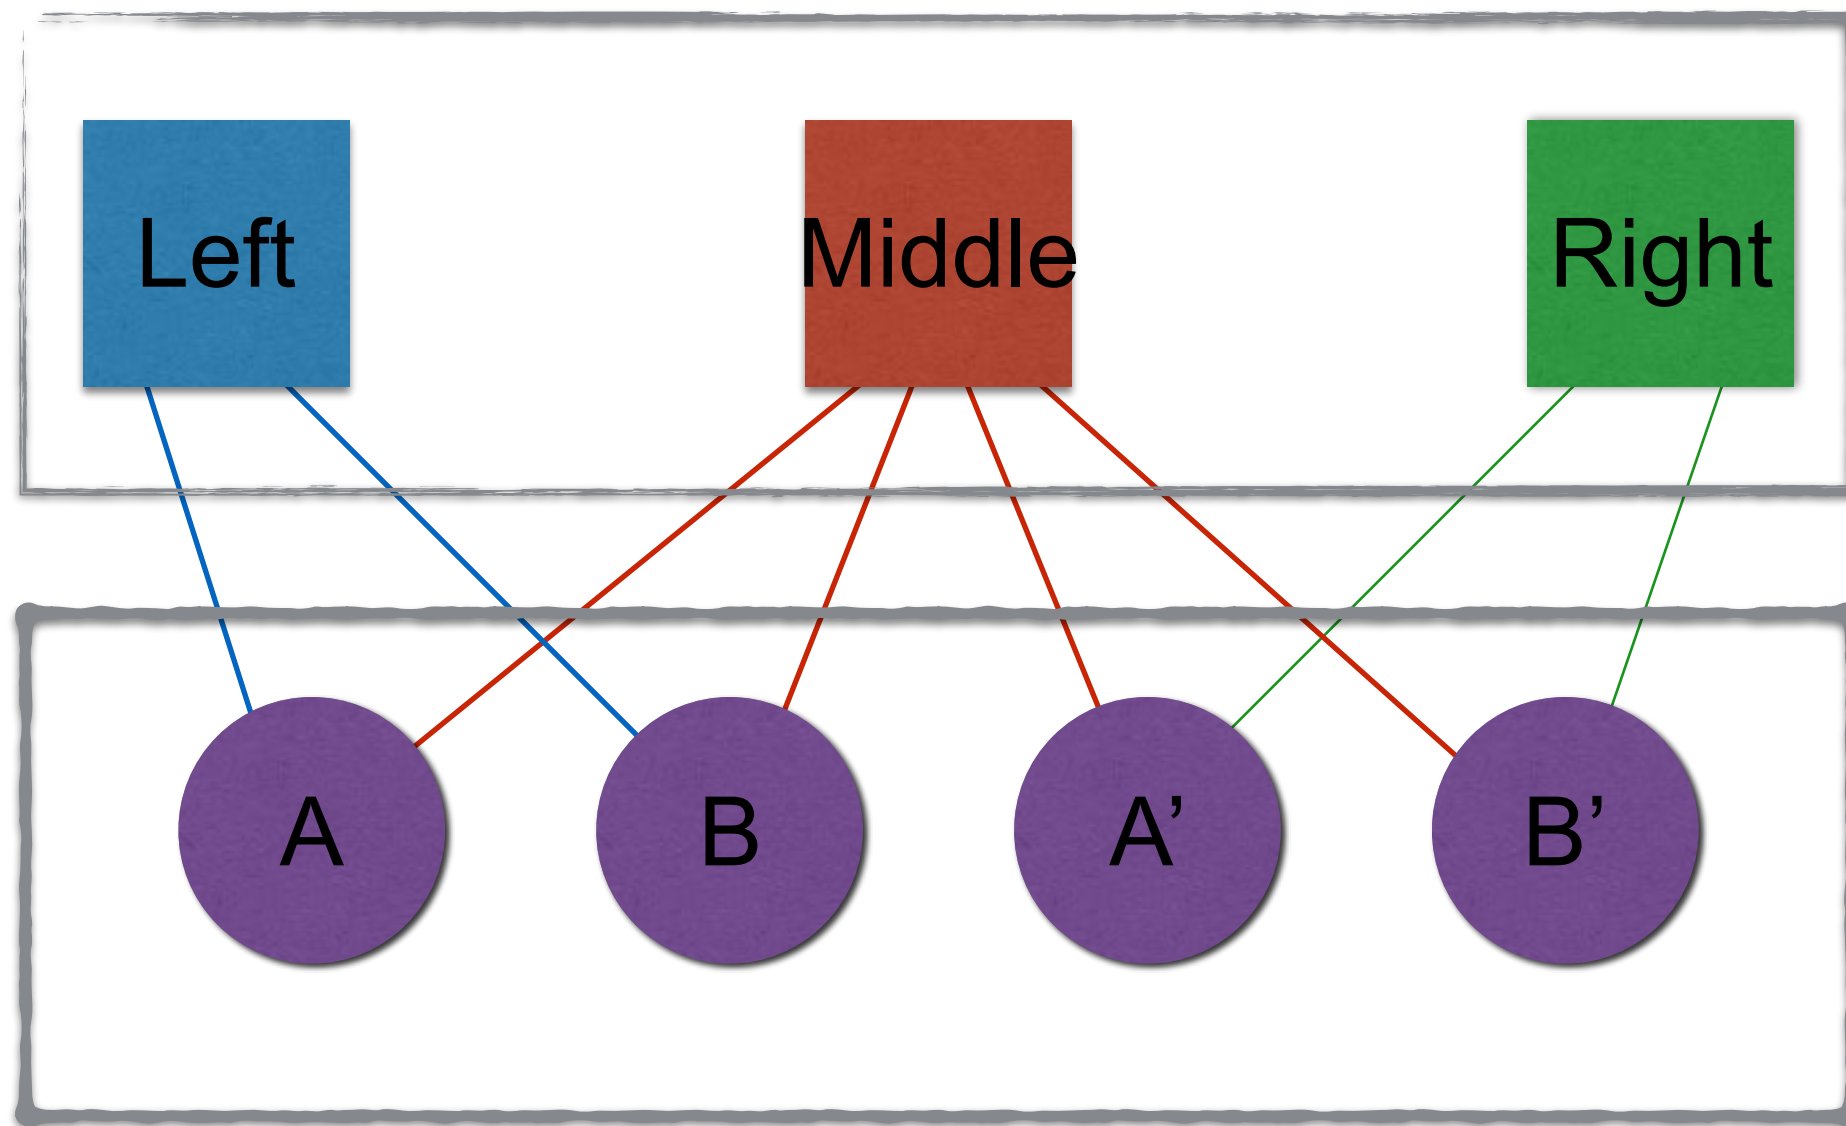
Networks



Akka cluster

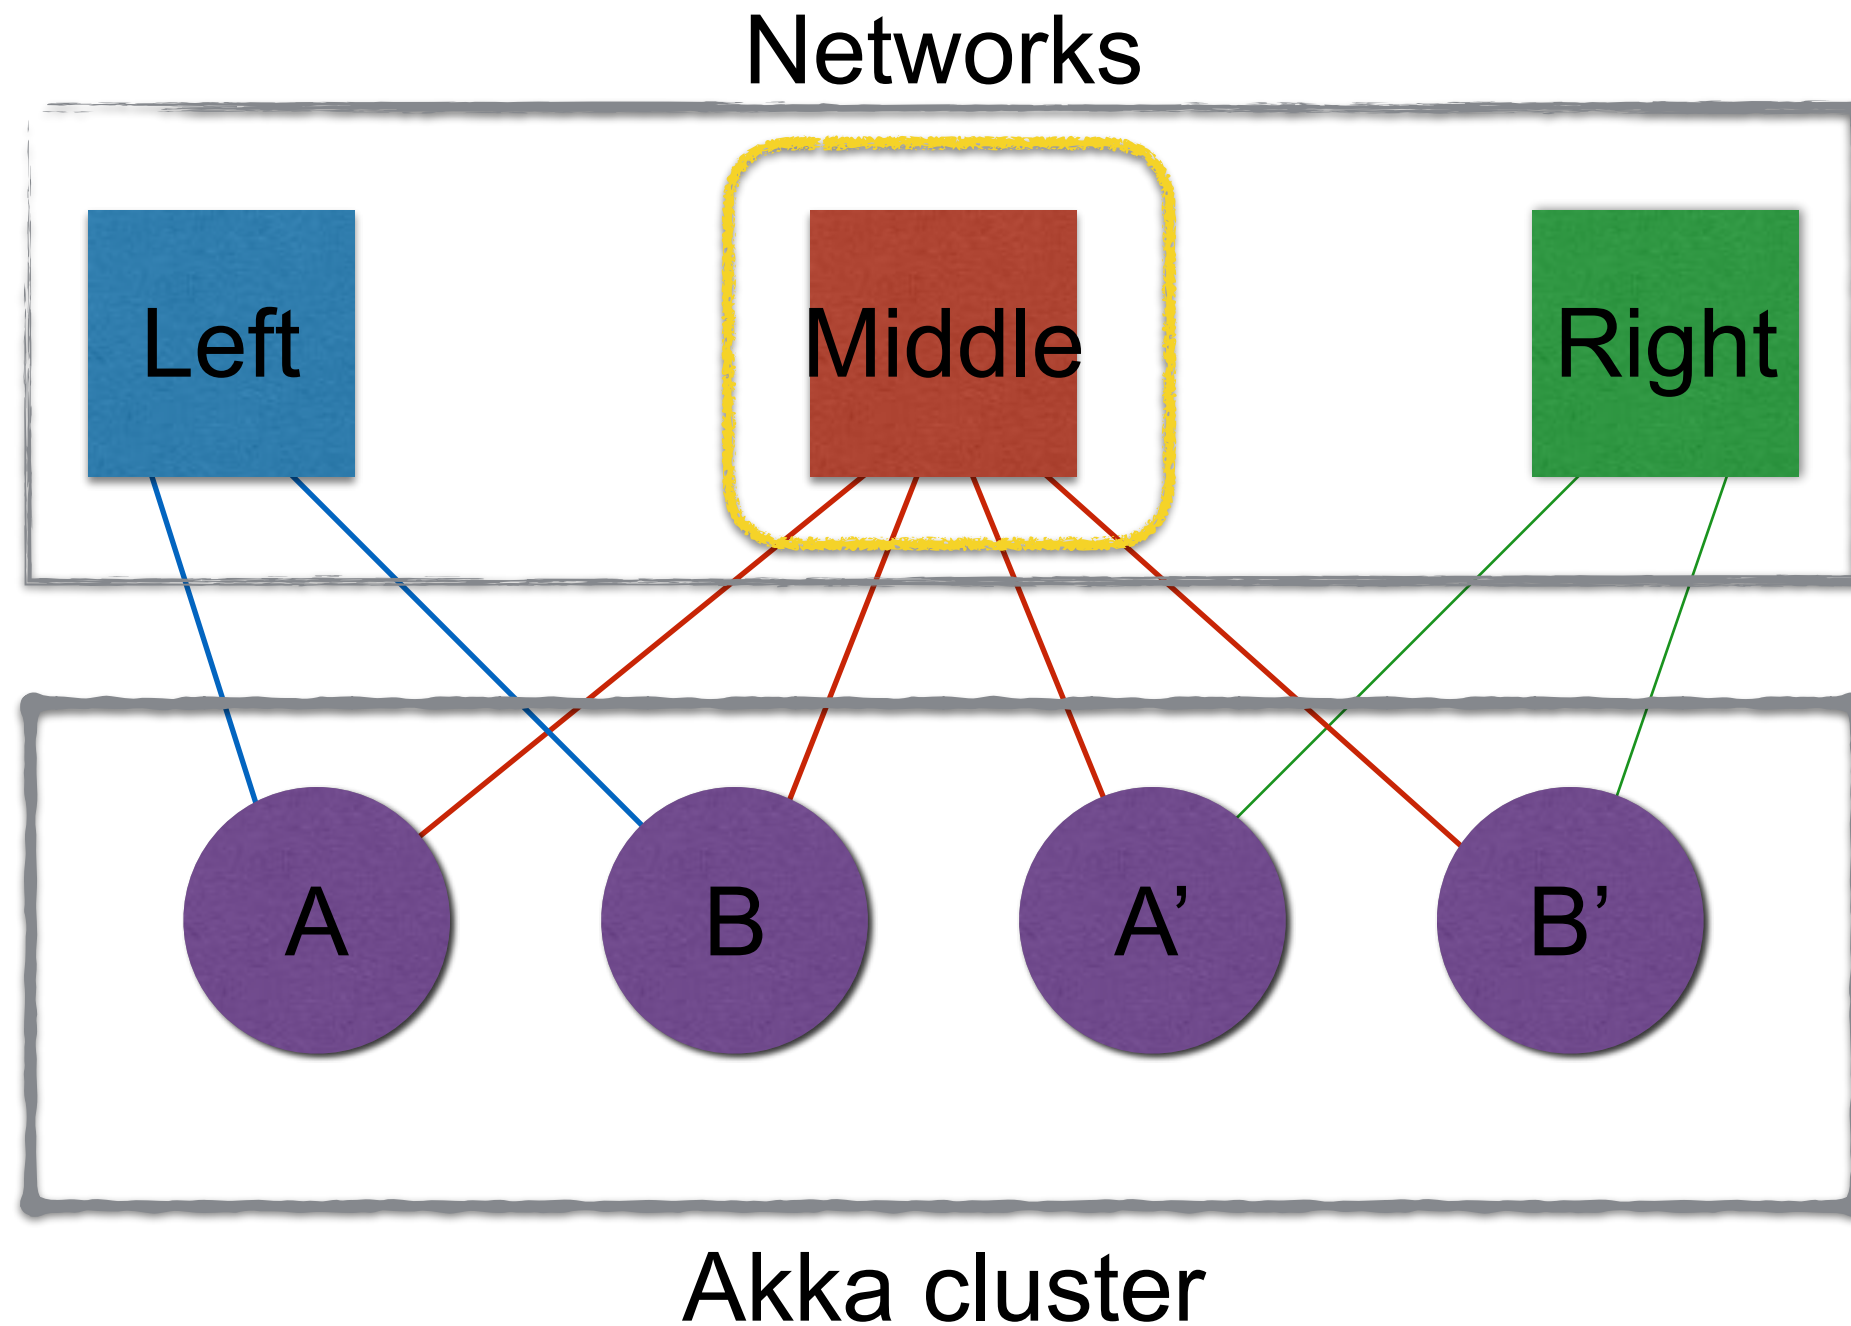
Demo Overview

Networks



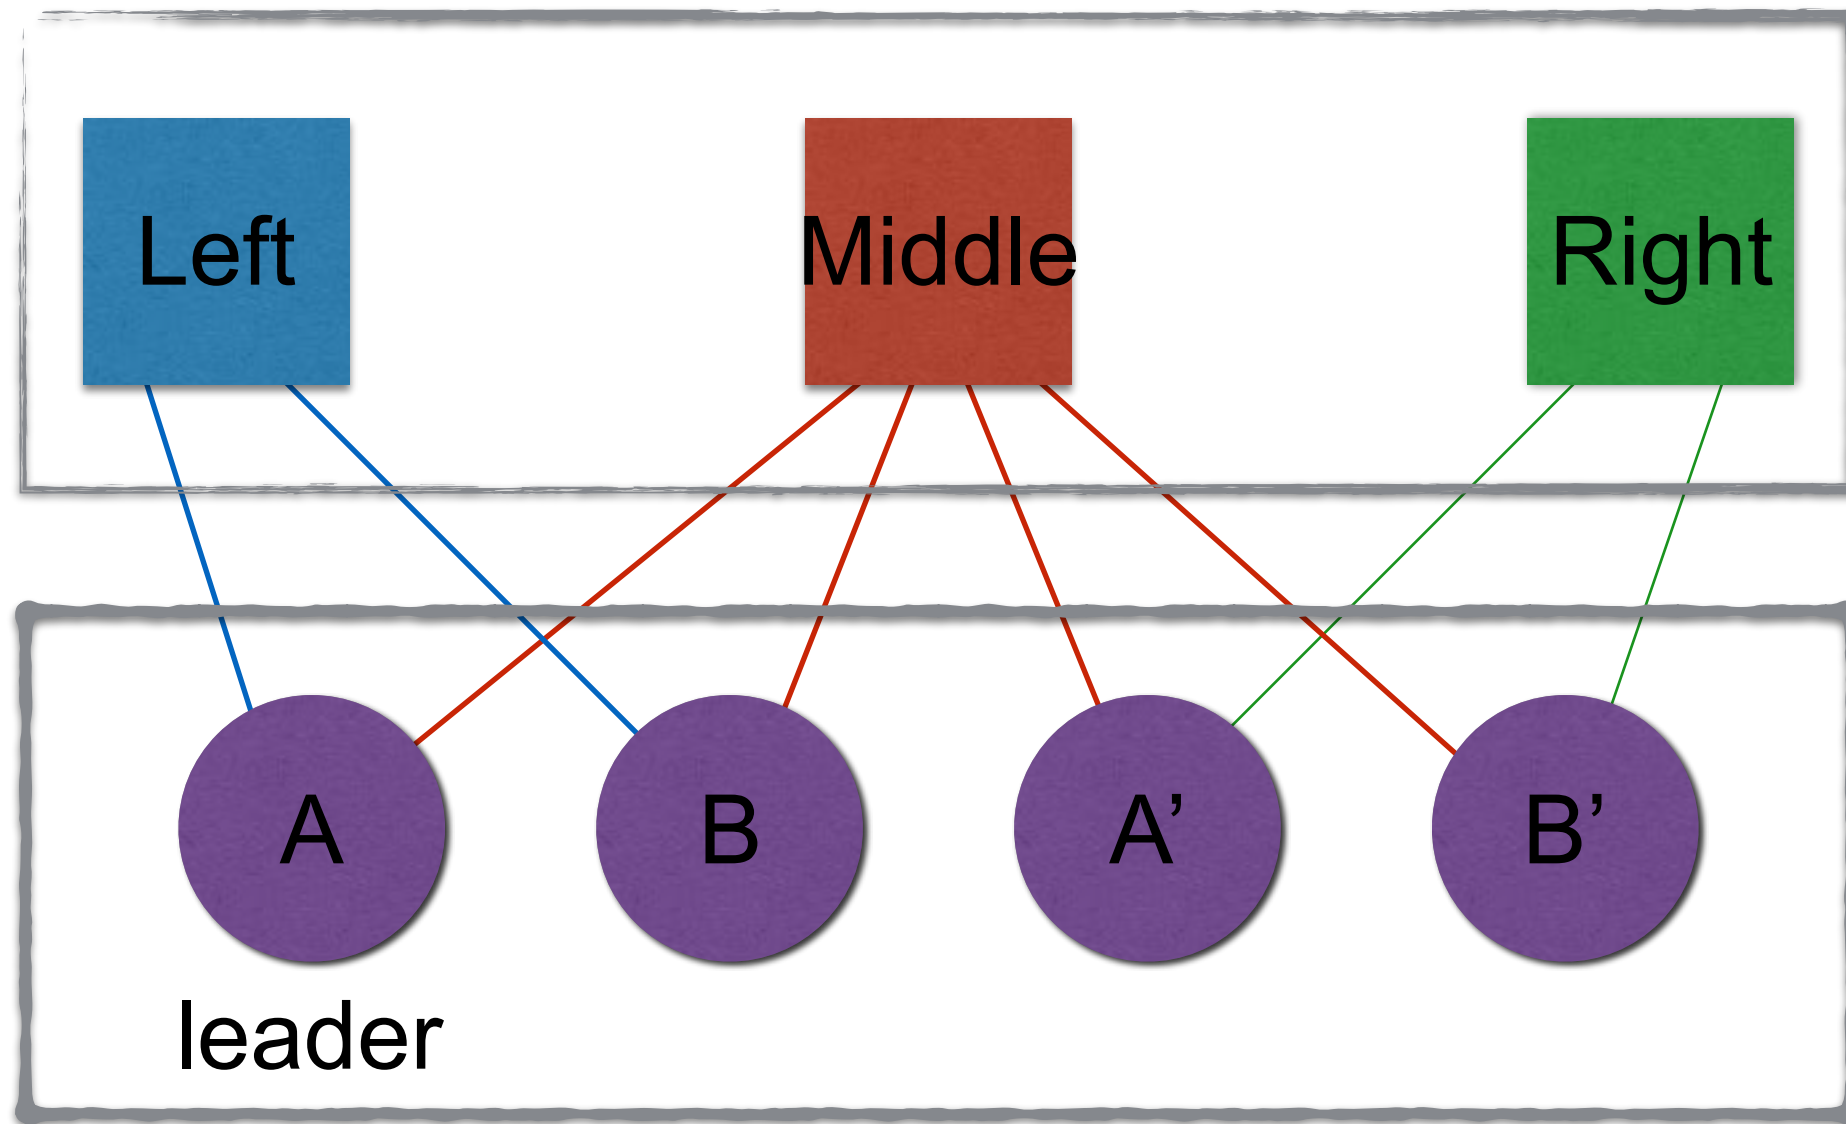
Akka cluster

Demo Overview



Demo Overview

Networks



Akka cluster

Auto-downing Enabled!!

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =
  s"""$name:
    template:
      resources:
        - cakesolutions.docker.jmx.akka
        - cakesolutions.docker.network.default.linux
      image: docker-compose-testkit-tests:$version
    environment:
      AKKA_HOST: $name
      AKKA_PORT: 2552
      CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"
    expose:
      - 2552
    networks:
      - $network1
      - $network2
  """.stripMargin

val yaml = DockerComposeString(
  s"""version: '2'

  services:
    ${clusterNode("A", "left", "middle")}
    ${clusterNode("B", "left", "middle")}
    ${clusterNode("A", "right", "middle")}
    ${clusterNode("B", "right", "middle")}

  networks:
    left:
    middle:
    right:
  """.stripMargin
)
```

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:  
    template:  
      resources:  
        - cakesolutions.docker.jmx.akka  
        - cakesolutions.docker.network.default.linux  
      image: docker-compose-testkit-tests:$version  
    environment:  
      AKKA_HOST: $name  
      AKKA_PORT: 2552  
      CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"  
    expose:  
      - 2552  
    networks:  
      - $network1  
      - $network2  
  """.stripMargin
```

```
val yaml = DockerComposeString(  
  s"""version: '2'  
  
  services:  
    ${clusterNode("A", "left", "middle")}  
    ${clusterNode("B", "left", "middle")}  
    ${clusterNode("A", "right", "middle")}  
    ${clusterNode("B", "right", "middle")}  
  
  networks:  
    left:  
    middle:  
    right:  
  """.stripMargin  
)
```

Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =
  s"""$name:
    template:
      resources:
        - cakesolutions.docker.jmx.akka
        - cakesolutions.docker.network.default.linux
      image: docker-compose-testkit-tests:$version
    environment:
      AKKA_HOST: $name
      AKKA_PORT: 2552
      CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"
    expose:
      - 2552
    networks:
      - $network1
      - $network2
  """.stripMargin
```

```
val yaml = DockerComposeString(
  s"""version: '2'

  services:
    ${clusterNode("A", "left", "middle")}
    ${clusterNode("B", "left", "middle")}
    ${clusterNode("A", "right", "middle")}
    ${clusterNode("B", "right", "middle")}

  networks:
    left:
    middle:
    right:
  """.stripMargin
)
```



Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =
  s"""$name:
    template:
      resources:
        - cakesolutions.docker.jmx.akka
        - cakesolutions.docker.network.default.linux
      image: docker-compose-testkit-tests:$version
      AKKA_HOST: $name
      AKKA_PORT: 2552
      CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"
    expose:
      - 2552
    networks:
      - $network1
      - $network2
  """.stripMargin

val yaml = DockerComposeString(
  s"""version: '2'

  services:
    ${clusterNode("A", "left", "middle")}
    ${clusterNode("B", "left", "middle")}
    ${clusterNode("A", "right", "middle")}
    ${clusterNode("B", "right", "middle")}

  networks:
    left:
    middle:
    right:
  """.stripMargin
)
```


Demo Overview

```
def clusterNode(name: String, network1: String, network2: String): String =  
  s"""$name:
```

```
    template:  
      resources:  
        - cakesolutions.docker.jmx.akka  
        - cakesolutions.docker.network.default.linux  
        image: docker-compose-testkit-tests:$version  
        AKKA_HOST: $name  
        AKKA_PORT: 2552  
        CLUSTER_SEED_NODE: "akka.tcp://TestCluster@A:2552"  
      expose:  
        - 2552  
      networks:  
        - $network1  
        - $network2  
  """.stripMargin
```

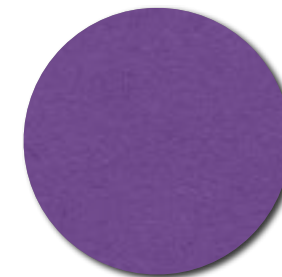
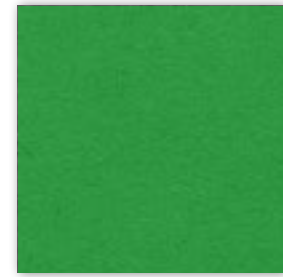
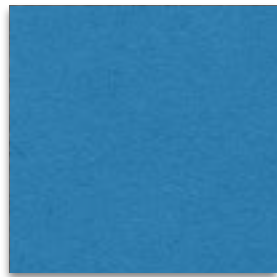
immutable container

```
val yaml = DockerComposeString(  
  s"""version: '2'
```

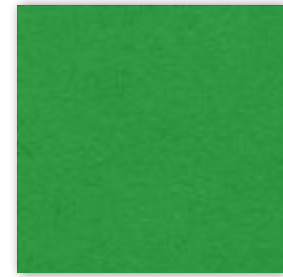
```
    services:  
      ${clusterNode("A", "left", "middle")}  
      ${clusterNode("B", "left", "middle")}  
      ${clusterNode("A", "right", "middle")}  
      ${clusterNode("B", "right", "middle")}  
    networks:  
      left:  
      middle:  
      right:  
  """.stripMargin  
)
```



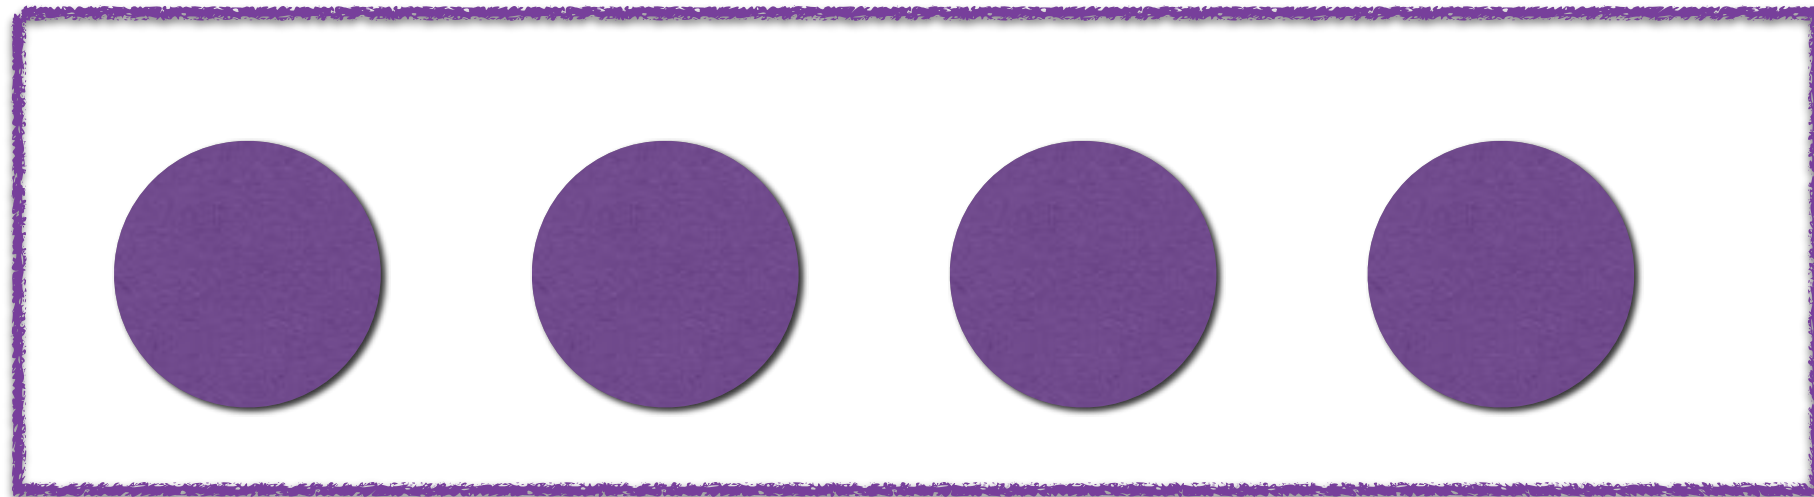
Demo Overview



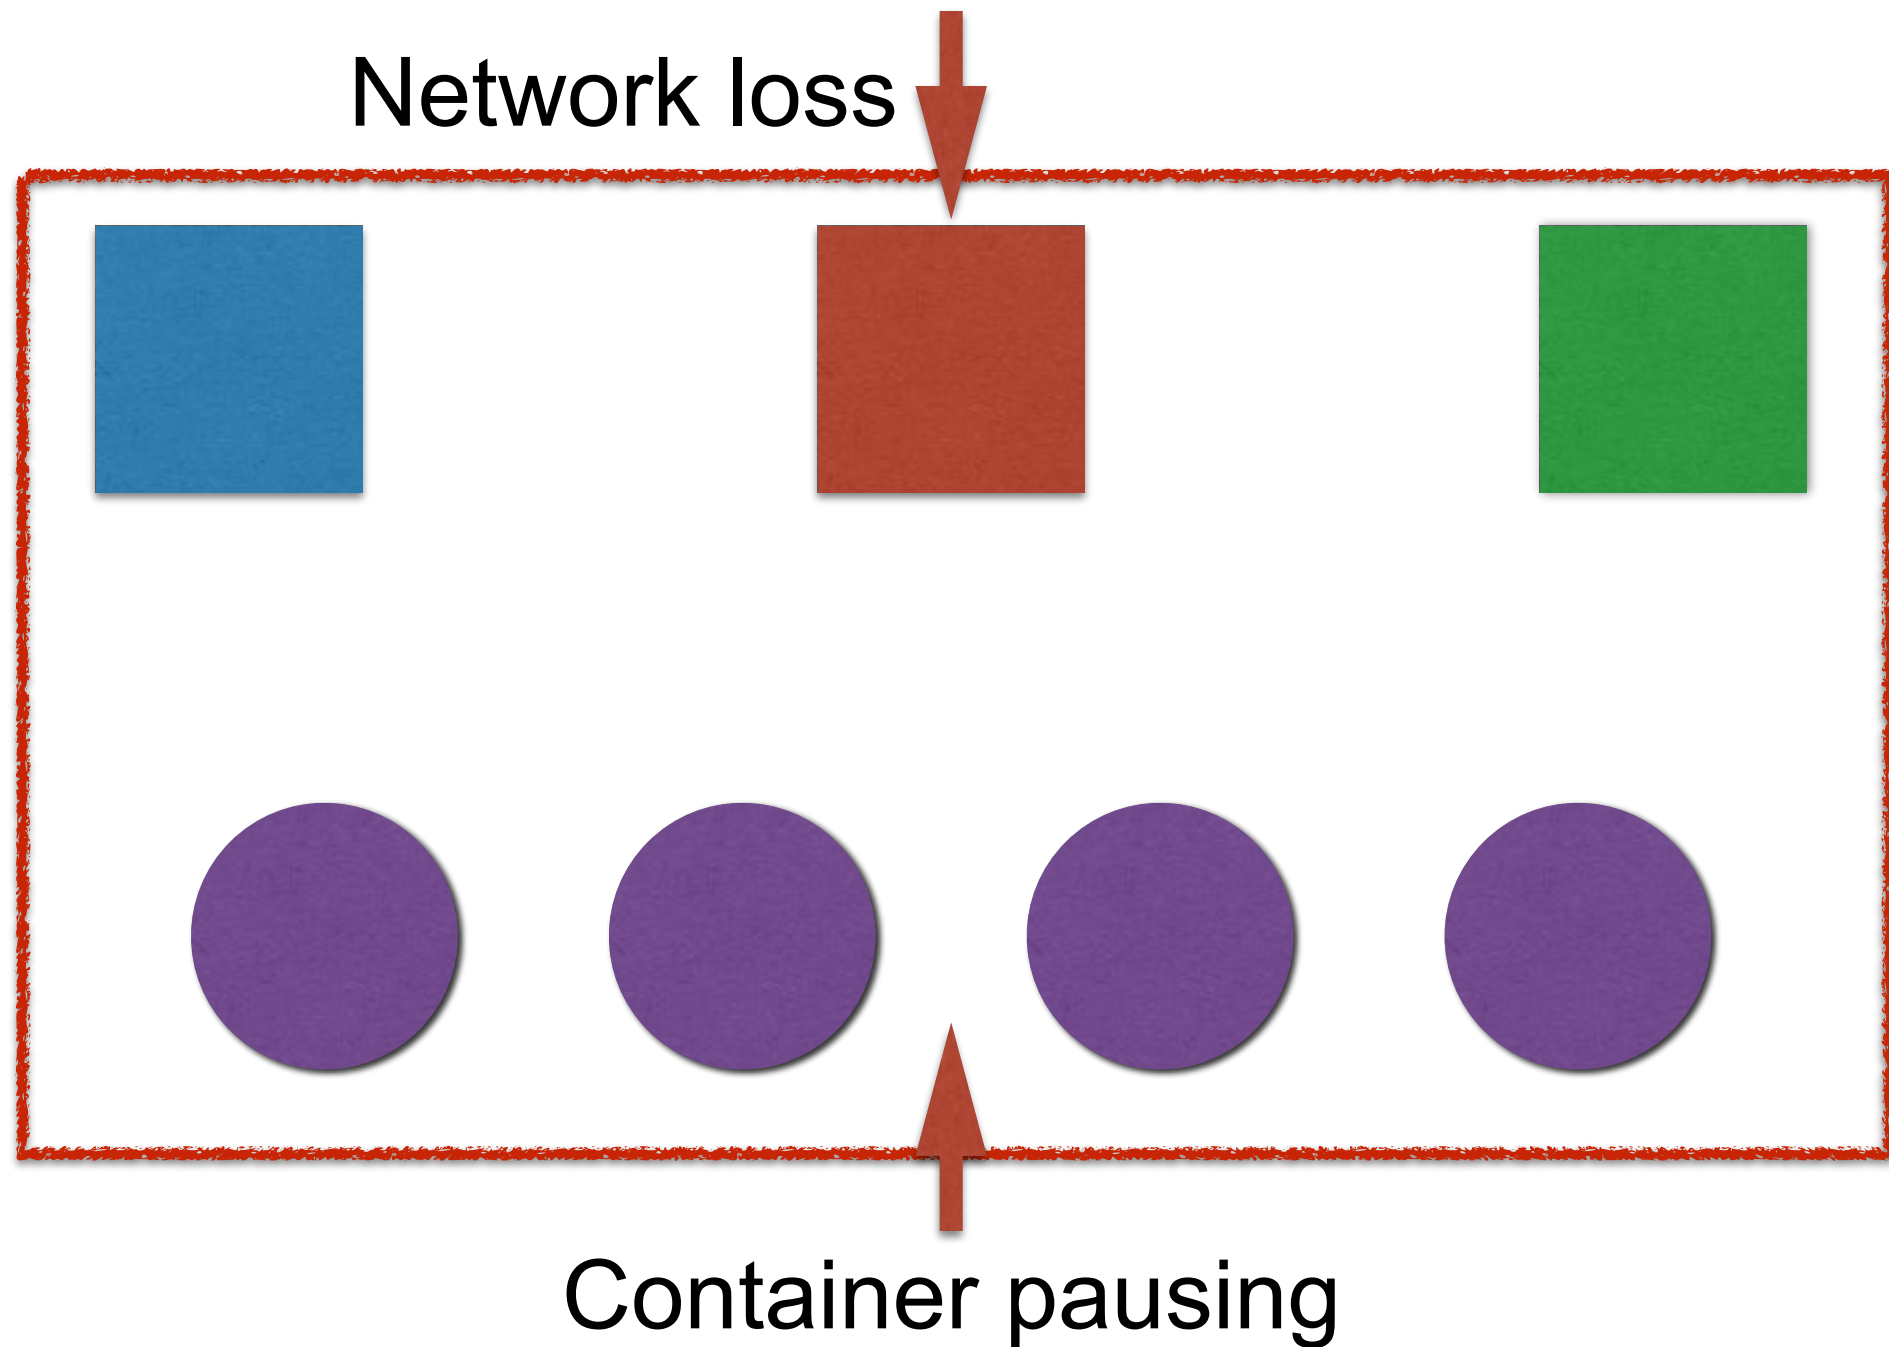
Demo Overview



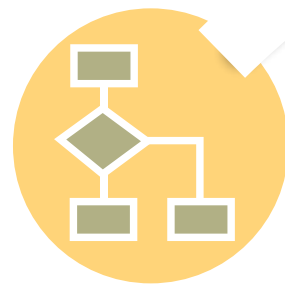
Akka JMX
Console



Demo Overview

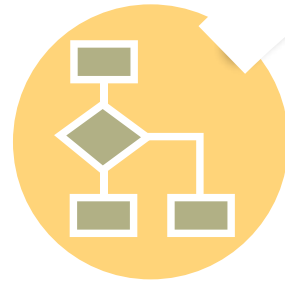


Chaos Experiment

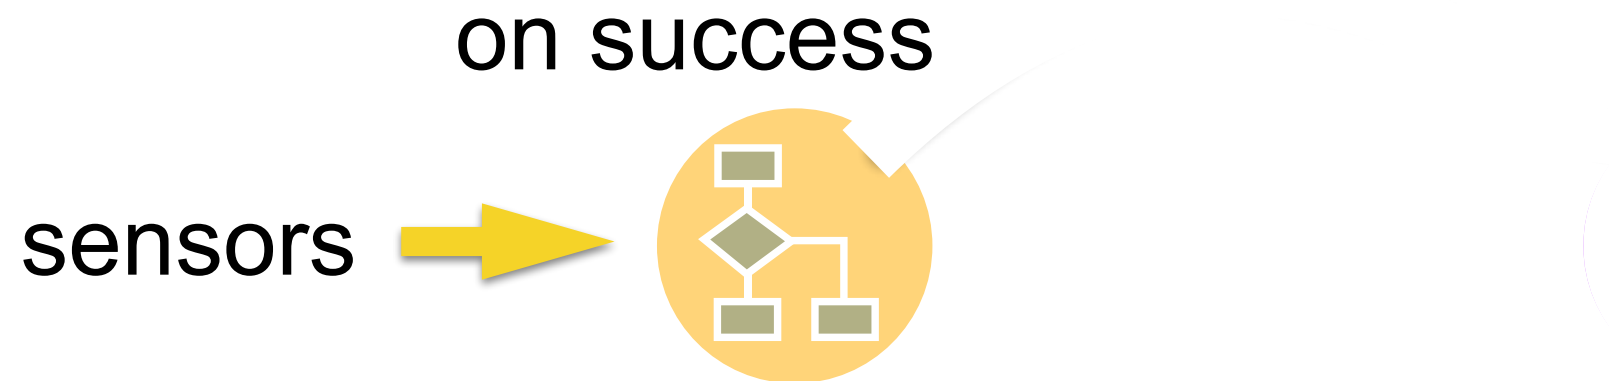


Chaos Experiment

sensors

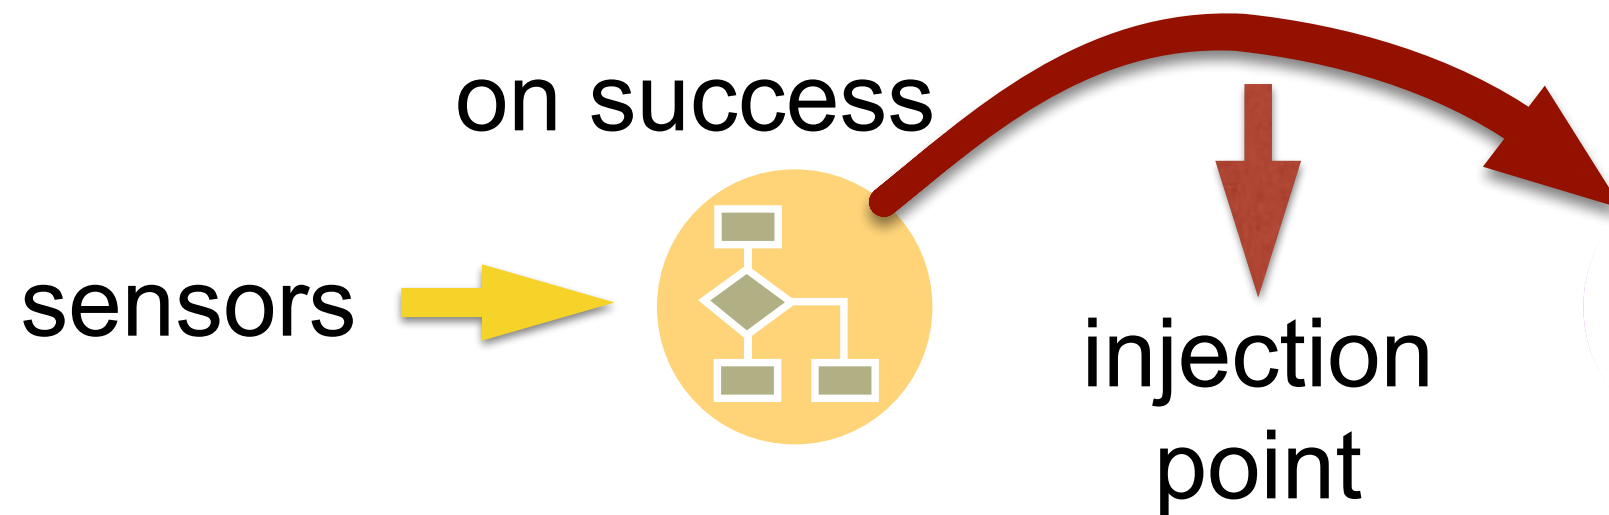


Chaos Experiment



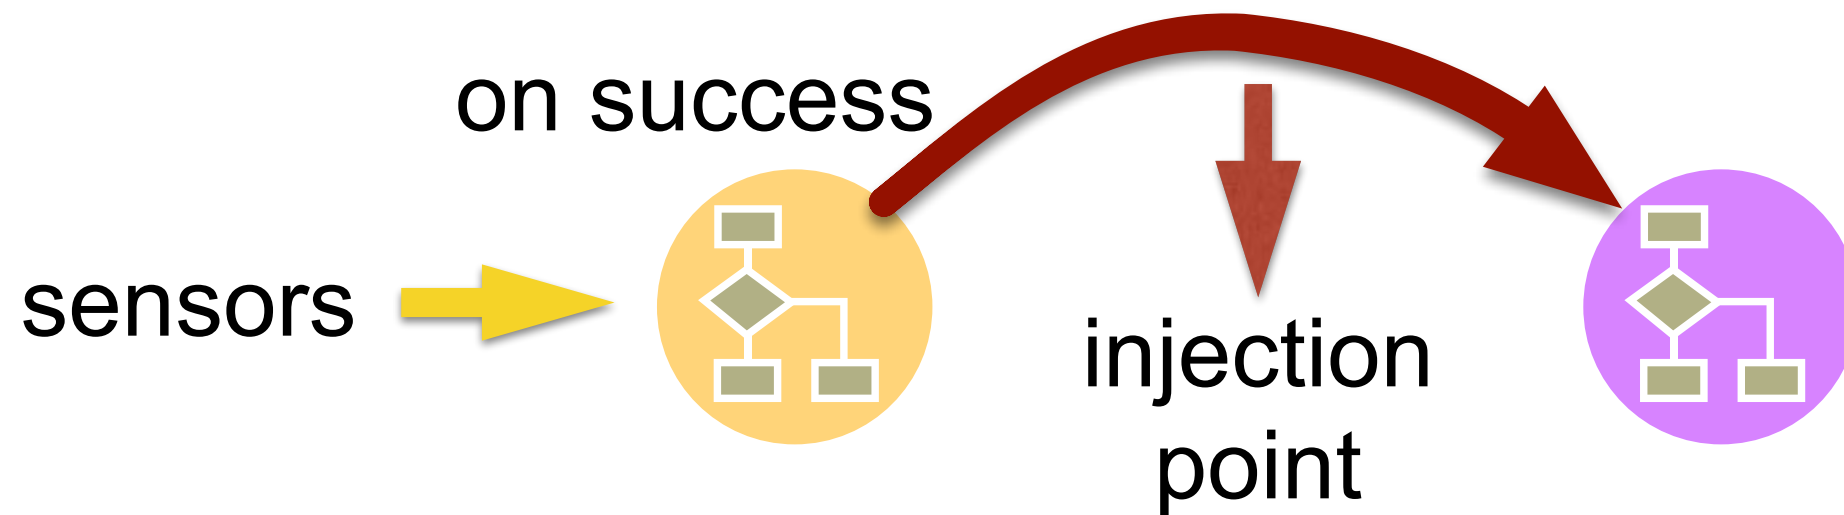
- Interpret this as follows:
 - when we successfully detect behaviour 

Chaos Experiment



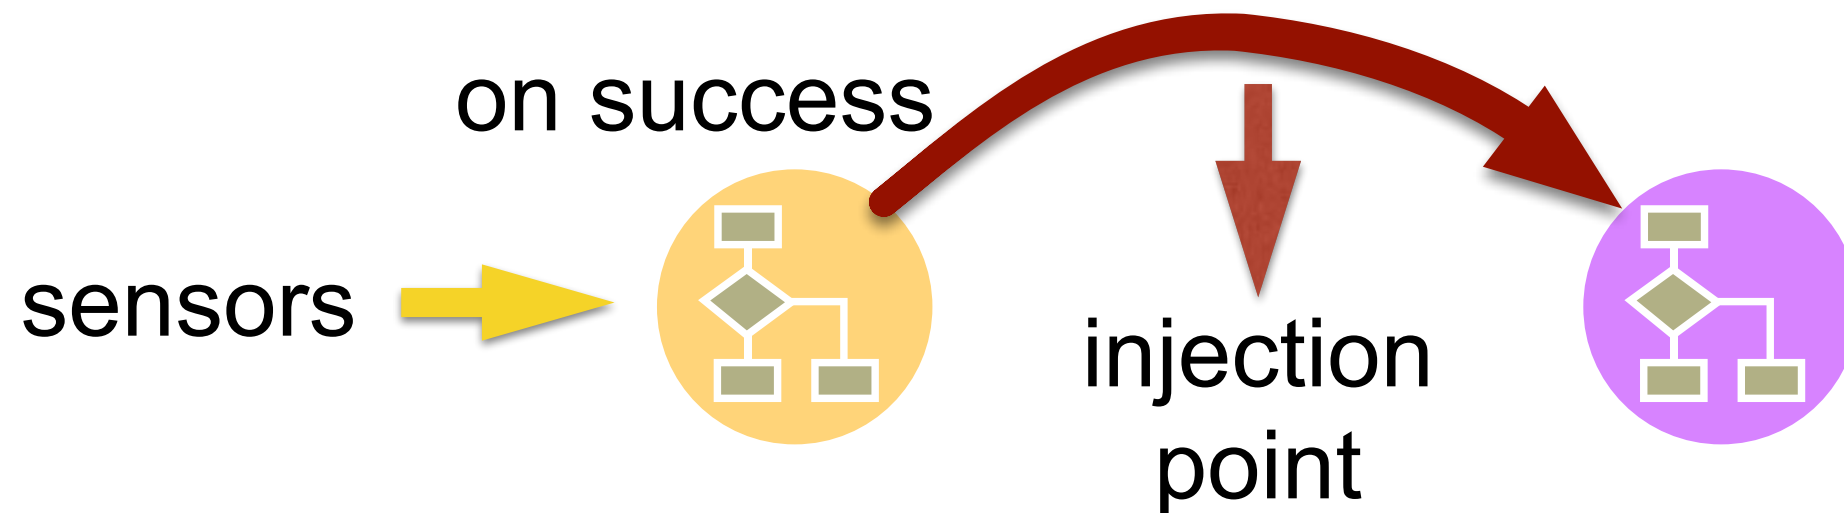
- Interpret this as follows:
 - when we successfully detect behaviour ●
 - action ➡ may happen

Chaos Experiment



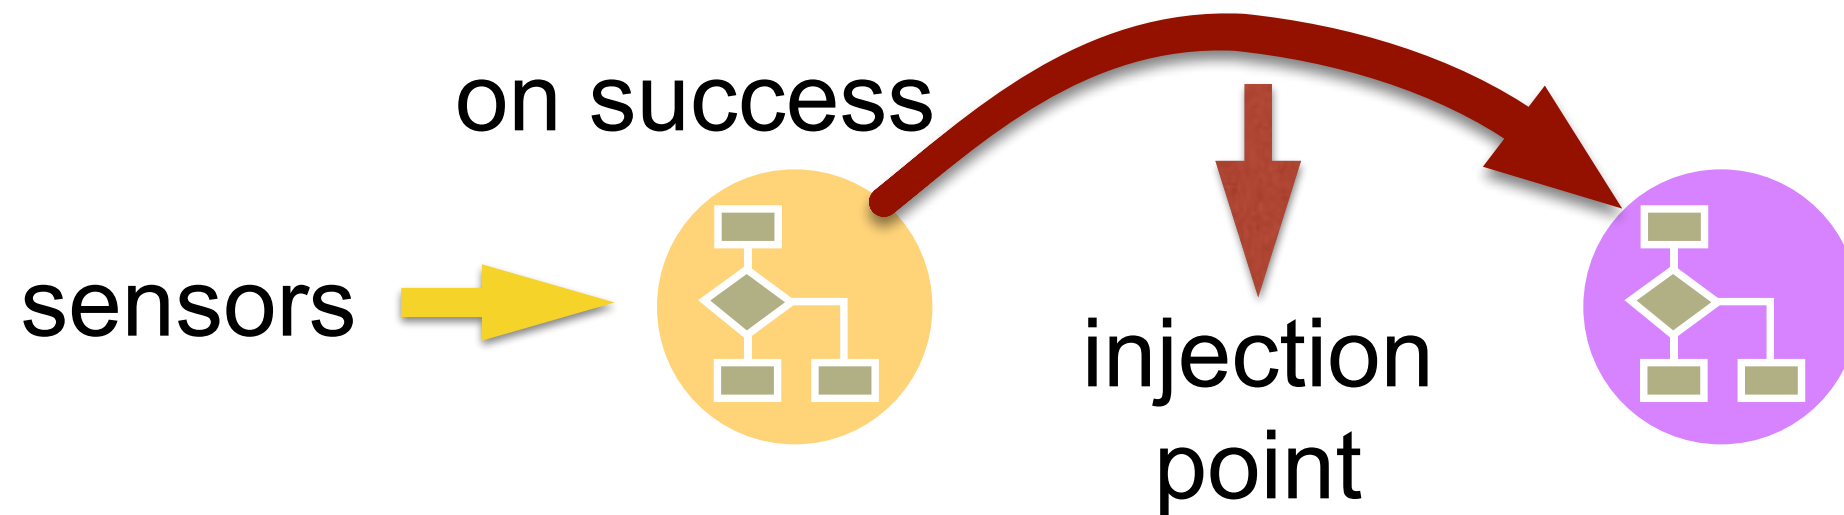
- Interpret this as follows:
 - when we successfully detect behaviour ●
 - action ➡ may happen
 - and we start monitoring for behaviour ●

Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour ●
 - action ➡ may happen
 - and we start monitoring for behaviour ●
- Chaining such relations allows the fault injection search space to be described

Chaos Experiment



- Interpret this as follows:
 - when we successfully detect behaviour ●
 - action ➡ may happen
 - and we start monitoring for behaviour ●
- Chaining such relations allows the fault injection search space to be described

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {
```

```
} yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- ???  
    ..  
  
  } yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- ???  
    ..  
    _ <- jvmGC(JvmGCStart)(rightA)  
    _ = note("A' JVM GC pause starts")  
    obs2 <- ???  
    ..  
  
  } yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- ???  
    ..  
    _ <- jvmGC(JvmGCStart)(rightA)  
    _ = note("A' JVM GC pause starts")  
    obs2 <- ???  
    ..  
    _ <- jvmGC(JvmGCEnd)(rightA)  
    _ = note("A' JVM GC pause ends")  
    obs3 <- ???  
    ..  
  } yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- ???  
    ..  
    _ <- jvmGC(JvmGCStart)(rightA)  
    _ = note("A' JVM GC pause starts")  
    obs2 <- ???  
    ..  
    _ <- jvmGC(JvmGCEnd)(rightA)  
    _ = note("A' JVM GC pause ends")  
    obs3 <- ???  
    ..  
    _ <- impair(Loss("100%"))("middle")  
    _ = note("partition into left and right networks")  
    obs4 <- ???  
    ..  
  } yield Accept()
```


Demo Overview

```
def inCluster(nodes: Address*): Monitor[Boolean, AkkaClusterState] = {  
  Monitor(false) {  
    case false => {  
      case Observe(AkkaClusterState(_, members, unreachable))  
        if unreachable.isEmpty && nodes.forall(node => members.filter(_.status ==  
Up).exists(_.address == node)) =>  
        Goto(true, 3.seconds)  
    }  
  
    case true => {  
      case Observe(AkkaClusterState(_, _, unreachable)) if unreachable.nonEmpty =>  
        Stop(Fail())  
      case Observe(AkkaClusterState(_, members, _))  
        if nodes.exists(node => members.filter(_.status == Up).forall(_.address !=  
node)) =>  
        Stop(Fail())  
      case StateTimeout =>  
        Stop(Accept())  
    }  
  }  
}
```

Demo Overview

Monitor State

```
def inCluster(nodes: Address*): Monitor[Boolean, AkkaClusterState] = {  
  Monitor(false) { Initial State  
    case false => {  
      case observe(AkkaClusterState(_, members, unreachable))  
        if unreachable.isEmpty && nodes.forall(node => members.filter(_.status ==  
Up).exists(_.address == node)) =>  
        Goto(true, 3.seconds)  
    }  
  
    case true => {  
      case observe(AkkaClusterState(_, _, unreachable)) if unreachable.nonEmpty =>  
        Stop(Fail())  
      case observe(AkkaClusterState(_, members, _))  
        if nodes.exists(node => members.filter(_.status == Up).forall(_.address !=  
node)) =>  
        Stop(Fail())  
      case StateTimeout =>  
        Stop(Accept())  
    }  
  }  
}
```



Demo Overview

Observed Event

```
def inCluster(nodes: Address*): Monitor[Boolean, AkkaClusterState] = {  
  Monitor(false) {  
    case false => {  
      case Observe(AkkaClusterState(_, members, unreachable))  
        if unreachable.isEmpty && nodes.forall(node => members.filter(_.status ==  
Up).exists(_.address == node)) =>  
        Goto(true, 3.seconds)  
    }  
  
    case true => {  
      case Observe(AkkaClusterState(_, _, unreachable)) if unreachable.nonEmpty =>  
        Stop(Fail())  
      case Observe(AkkaClusterState(_, members, _))  
        if nodes.exists(node => members.filter(_.status == Up).forall(_.address !=  
node)) =>  
        Stop(Fail())  
      case StateTimeout =>  
        Stop(Accept())  
    }  
  }  
}
```



Demo Overview

```
def inCluster(nodes: Address*): Monitor[Boolean, AkkaClusterState] = {  
  Monitor(false) {  
    case false => {  
      case Observe(AkkaClusterState(_, members, unreachable))  
        if unreachable.isEmpty && nodes.forall(node => members.filter(_.status ==  
Up).exists(_.address == node)) =>  
        Goto(true, 3.seconds)  
    }  
  
    case true => {  
      case Observe(AkkaClusterState(_, _, unreachable)) if unreachable.nonEmpty =>  
        Stop(Fail())  
      case Observe(AkkaClusterState(_, members, _))  
        if nodes.exists(node => members.filter(_.status == Up).forall(_.address !=  
node)) =>  
        Stop(Fail())  
      case StateTimeout => Timeout Event  
        Stop(Accept())  
    }  
  }  
}
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {
```

```
} yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs1))  
    _ = note("cluster formed, is stable and A an available leader")  
    ..  
  
  } yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs1))  
    _ = note("cluster formed, is stable and A an available leader")  
    ..  
    obs2 <- jmx(unreachable(rightA))(leftA)  
    _ <- check(isAccepting(obs2))  
    _ = note("A' is unreachable")  
    ..  
  
  } yield Accept()
```

Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs1))  
    _ = note("cluster formed, is stable and A an available leader")  
    ..  
    obs2 <- jmx(unreachable(rightA))(leftA)  
    _ <- check(isAccepting(obs2))  
    _ = note("A' is unreachable")  
    ..  
    obs3 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs3))  
    _ = note("cluster stabilised with A' as a member")  
    ..  
  } yield Accept()
```


Demo Overview

```
def experiment[M: _jvm: _network: ..](implicit ..): Eff[M, Notify] =  
  for {  
    obs1 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs1))  
    _ = note("cluster formed, is stable and A an available leader")  
    ..  
    obs2 <- jmx(unreachable(rightA))(leftA)  
    _ <- check(isAccepting(obs2))  
    _ = note("A' is unreachable")  
    ..  
    obs3 <- jmx(inCluster(leftA, leftB, rightA, rightB))(leftA)  
    _ <- check(isAccepting(obs3))  
    _ = note("cluster stabilised with A' as a member")  
    ..  
    obs4 <- jmx(inCluster(leftA, leftB))(leftA)  
              && jmx(inCluster(leftA, leftB))(leftB)  
              && jmx(inCluster(rightA))(rightA)  
              && jmx(inCluster(rightB))(rightB)  
    _ <- check(isAccepting(obs4))  
    _ = note("cluster split brains into 3 clusters: A & B; A'; B'")  
  } yield Accept()
```



Demo Overview

Demo Overview

```
type Model = Fx.fx5[JmxAction, NetworkAction, ..]
```

Demo Overview

```
val cluster = Map(  
  leftA -> compose.service(leftA).docker.head,  
  leftB -> compose.service(leftB).docker.head,  
  rightA -> compose.service(rightA).docker.head,  
  rightB -> compose.service(rightB).docker.head  
)
```

Demo Overview

```
experiment[Model].runJvm(cluster).runNetwork...
```

Demo Overview

```
experiment[Model].runJvm(cluster).runNetwork...
```



CAKESOLUTIONS

Demo Time!



MANCHESTER

LONDON

NEW YORK

Conclusion

Conclusion

- Seen how to model a variety of distributed systems extensionally

Conclusion

- Seen how to model a variety of distributed systems extensionally
- Taken steady state behavioural specifications and developed a composable DSL for experiments
 - by relating data injection traces to observed system behaviour

Conclusion

- Seen how to model a variety of distributed systems extensionally
- Taken steady state behavioural specifications and developed a composable DSL for experiments
 - by relating data injection traces to observed system behaviour
- Applied techniques to a range of fault-injection scenarios

Next Steps

Next Steps

- Prove out library on a range of real-world examples
 - avoid closed work assumptions

Next Steps

- Prove out library on a range of real-world examples
 - avoid closed work assumptions
- Describe how we may automatically generate Chaos experiments from models
 - observational models
 - models based on static analysis
 - models built using domain expertise

References

- <https://github.com/carlpulley/docker-compose-testkit/tree/scala-meetup-krakow-2017>
- *Chaos Engineering* by A.Basiri, N.Behnam, R.de Rooij, L.Hochstein, L.Kosewski, J.Reynolds and C.Rosenthal
- <http://principlesofchaos.org/>
- *The Weird Machines in Proof-Carrying Code* by J.Vanegue
- *The Byzantine Generals Problem* by L.Lamport, R.Shostak and M.Pease
- *Impossibility of Distributed Consensus with One Faulty Process* by M.Fischer, N.Lynch and M.Paterson
- *Freer monads, more extensible effects* by O.Kiselyov and H.Ishii



CAKESOLUTIONS



@cakesolutions



0845 617 1200



enquiries@cakesolutions.net



MANCHESTER

LONDON

NEW YORK